

CLOUD COMPUTING

INTRODUCTION TO DATA SCIENCE

TIM KRASKA



CLICKER

How was the celebration of knowledge?

- A) Very Easy**
- B) Just ok**
- C) Tough**
- D) Spring break made me forget about it**
- E) I do not want to talk about it**

BACKGROUND OF CLOUD COMPUTING

1980's and 1990's: 52% growth in performance per year!


2002: The thermal wall

- Speed (frequency) peaks, but transistors keep shrinking

2000's: Multicore revolution

- 15-20 years later than predicted, we have hit the performance wall

2010's: Rise of Big Data



SOURCES DRIVING BIG DATA

It's All Happening On-line



Every:
Click

User Generated (Web & Mobile)

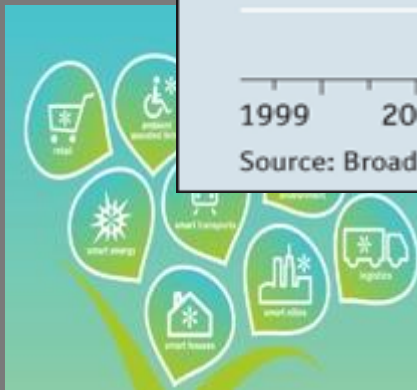


...

...




Internet of



Computing

Baseline information


Cost of genome sequencing compared with
Moore's law for computers



Source: Broad Institute

Baseline information

Cost of genome sequencing compared with
Moore's law for computers



Source: Broad Institute


DATA DELUGE

Billions of users connected through the net



- WWW, FB, twitter, cell phones, ...
- 80% of the data on FB was produced last year

Storage getting cheaper

- Store more data!



DATA GROWS FASTER THAN MOORE'S LAW



SOLVING THE IMPEDANCE MISMATCH

Computers not getting faster, and we are drowning in data

- How to resolve the dilemma?

Solution adopted by web-scale companies

- Go massively *distributed* and *parallel*



ENTER THE WORLD OF DISTRIBUTED SYSTEMS

Distributed Systems/Computing

- *Loosely coupled* set of computers, communicating through *message passing*, solving a common goal
- Tools: Msg passing, Distributed shared memory, RPC

Distributed computing is *challenging*

- Dealing with *partial failures* (examples?)
- Dealing with *asynchrony* (examples?)
- Dealing with *scale* (examples?)
- Dealing with *consistency* (examples?)


Distributed Computing versus Parallel Computing?

- distributed computing = parallel computing + partial failures

THE DATACENTER IS THE NEW COMPUTER

“The datacenter as a computer” still in its infancy

- Special purpose clusters, e.g., Hadoop cluster
- Built from less reliable components
- Highly variable performance
- Complex concepts are hard to program (low-level primitives)



= ?



DATA CENTER





DATACENTER/CLOUD COMPUTING OS

If the datacenter/cloud is the new computer

- What is its **Operating System**?
- Note that we are not talking about a host OS

Could be equivalent in benefit as the LAMP stack was to the .com boom – every startup *secretly* implementing the same functionality!

Open source stack for a Web 2.0 company:

- Linux OS
- Apache web server
- MySQL, MariaDB or MongoDB DBMS
- PHP, Perl, or Python languages for dynamic web pages

CLASSICAL OPERATING SYSTEMS

Data sharing

- Inter-Process Communication, RPC, files, pipes, ...

Programming Abstractions

- Libraries (libc), system calls, ...

Multiplexing of resources

- Scheduling, virtual memory, file allocation/protection, ...

DATACENTER/CLOUD OPERATING SYSTEM

Data sharing

- Google File System, [key/value stores](#)
- Apache project: Hadoop Distributed File System

Programming Abstractions

- Google MapReduce
- Apache projects: Hadoop, Pig, Hive, Spark

Multiplexing of resources

- Apache projects: Mesos, [YARN \(MapReduce v2\)](#), [ZooKeeper](#), [BookKeeper](#), ...

GOOGLE CLOUD INFRASTRUCTURE

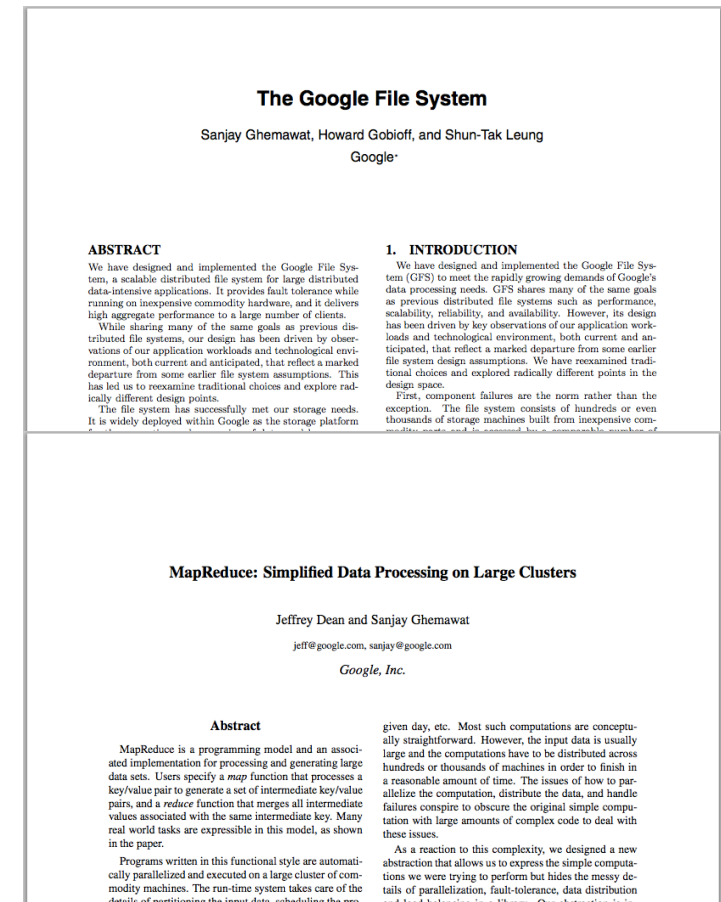
Google File System (GFS), 2003

- Distributed File System for entire cluster
- Single namespace

Google MapReduce (MR), 2004

- Runs queries/jobs on data
- Manages work distribution & fault-tolerance
- Colocated with file system

Apache open source versions: Hadoop DFS and Hadoop MR



HADOOP DISTRIBUTED FILE SYSTEM

Files split into 128MB *blocks*

Blocks replicated across several *datanodes* (usually 3)


Single *namenode* stores metadata (file names, block locations, etc)

Optimized for large files, sequential reads

Files are append-only

Data *striped* on hundreds/thousands of servers

- Scan 100 TB on 1 node @ 50 MB/s = 24 days
- Scan on 1000-node cluster = 35 minutes



CLICKER

The chance of a machine failing in 24h is 0.1%

What is the likelihood that one machine in a cluster of 1000 machines fails in 24h?

- a) 0.1%
- b) 10%
- c) 63%
- d) 99.999%

GFS/HDFS INSIGHTS (2)

***Failures* will be the norm**

Mean time between failures for 1 node = 3 years

Mean time between failures for 1000 nodes = **1 day**

Use *commodity* hardware

Failures are the norm anyway, buy cheaper hardware

No complicated consistency models

Single writer, append-only data

WHAT IS MAPREDUCE?

Simple data-parallel programming model designed for scalability and fault-tolerance

Pioneered by Google

- Processes 20 petabytes of data per day

Popularized by open-source Hadoop project

- Used at Yahoo!, Facebook, Amazon, ...



WHAT IS MAPREDUCE USED FOR?

- **At Google:**

- Index building for Google Search
- Article clustering for Google News
- Statistical machine translation


- **At Yahoo!:**

- Index building for Yahoo! Search
- Spam detection for Yahoo! Mail

- **At Facebook:**

- Data mining
- Ad optimization
- Spam detection

TYPICAL HADOOP CLUSTER



40 nodes/rack, 1000-4000 nodes in cluster

1 Gbps bandwidth within rack, 8 Gbps out of rack

Node specs (Yahoo terasort):

8 x 2GHz cores, 8 GB RAM, 4 disks (= 4 TB?)

CHALLENGES

Cheap nodes fail, especially if you have many

Mean time between failures for 1 node = 3 years

Mean time between failures for 1000 nodes = 1 day

Solution: Build fault-tolerance into system

Commodity network = low bandwidth

Solution: Push computation to the data

Programming distributed systems is hard

Solution: Data-parallel programming model: users write “map” & “reduce” functions, system distributes work and handles faults

MAPREDUCE PROGRAMMING MODEL

Data type: key-value *records*

Map function:

$$(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$$

Reduce function:

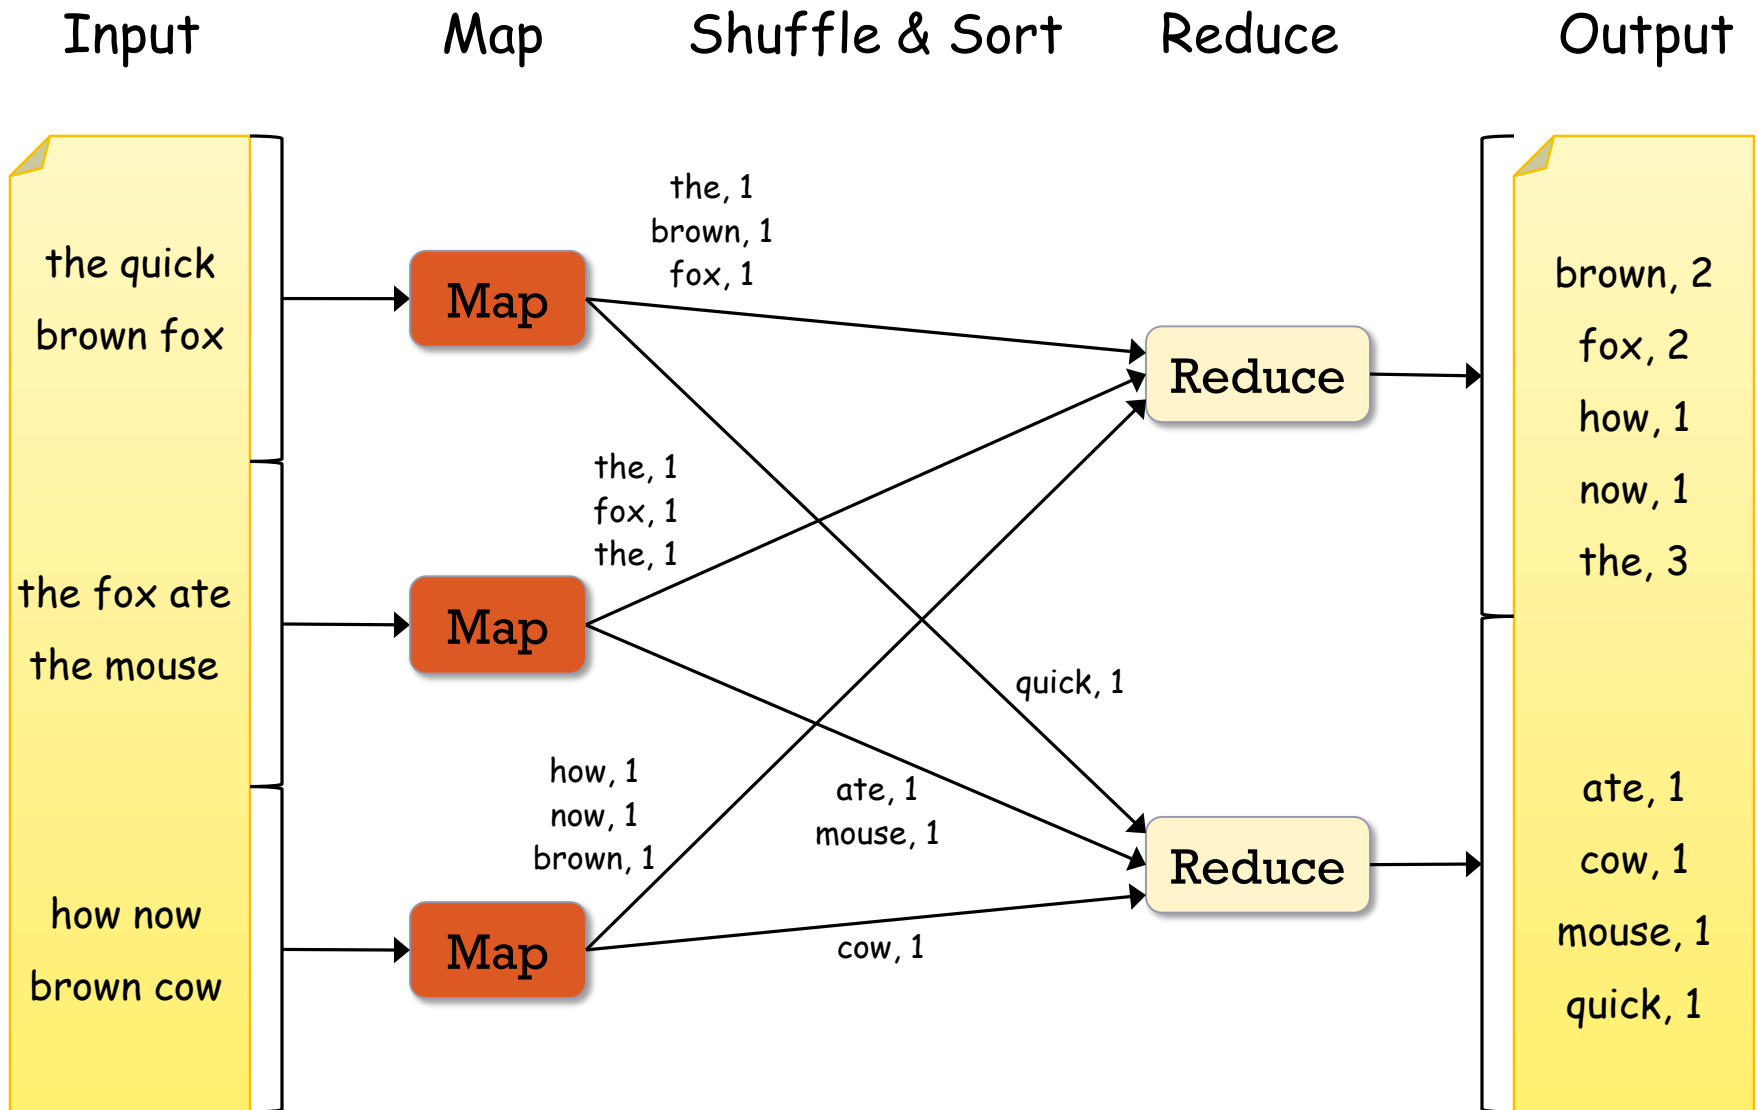
$$(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$$

EXAMPLE: WORD COUNT

```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reducer(key, values):  
    output(key, sum(values))
```

WORD COUNT EXECUTION



MAPREDUCE EXECUTION DETAILS

Single *master* controls job execution on multiple *slaves*

Mappers preferentially placed on same node or same rack as their input block

- Minimizes network usage

Mappers save outputs to local disk before serving them to reducers

- Allows recovery if a reducer crashes
- Allows having more reducers than nodes

AN OPTIMIZATION: THE COMBINER

A combiner is a local aggregation function for repeated keys produced by same map

Works for associative functions like sum, count, max

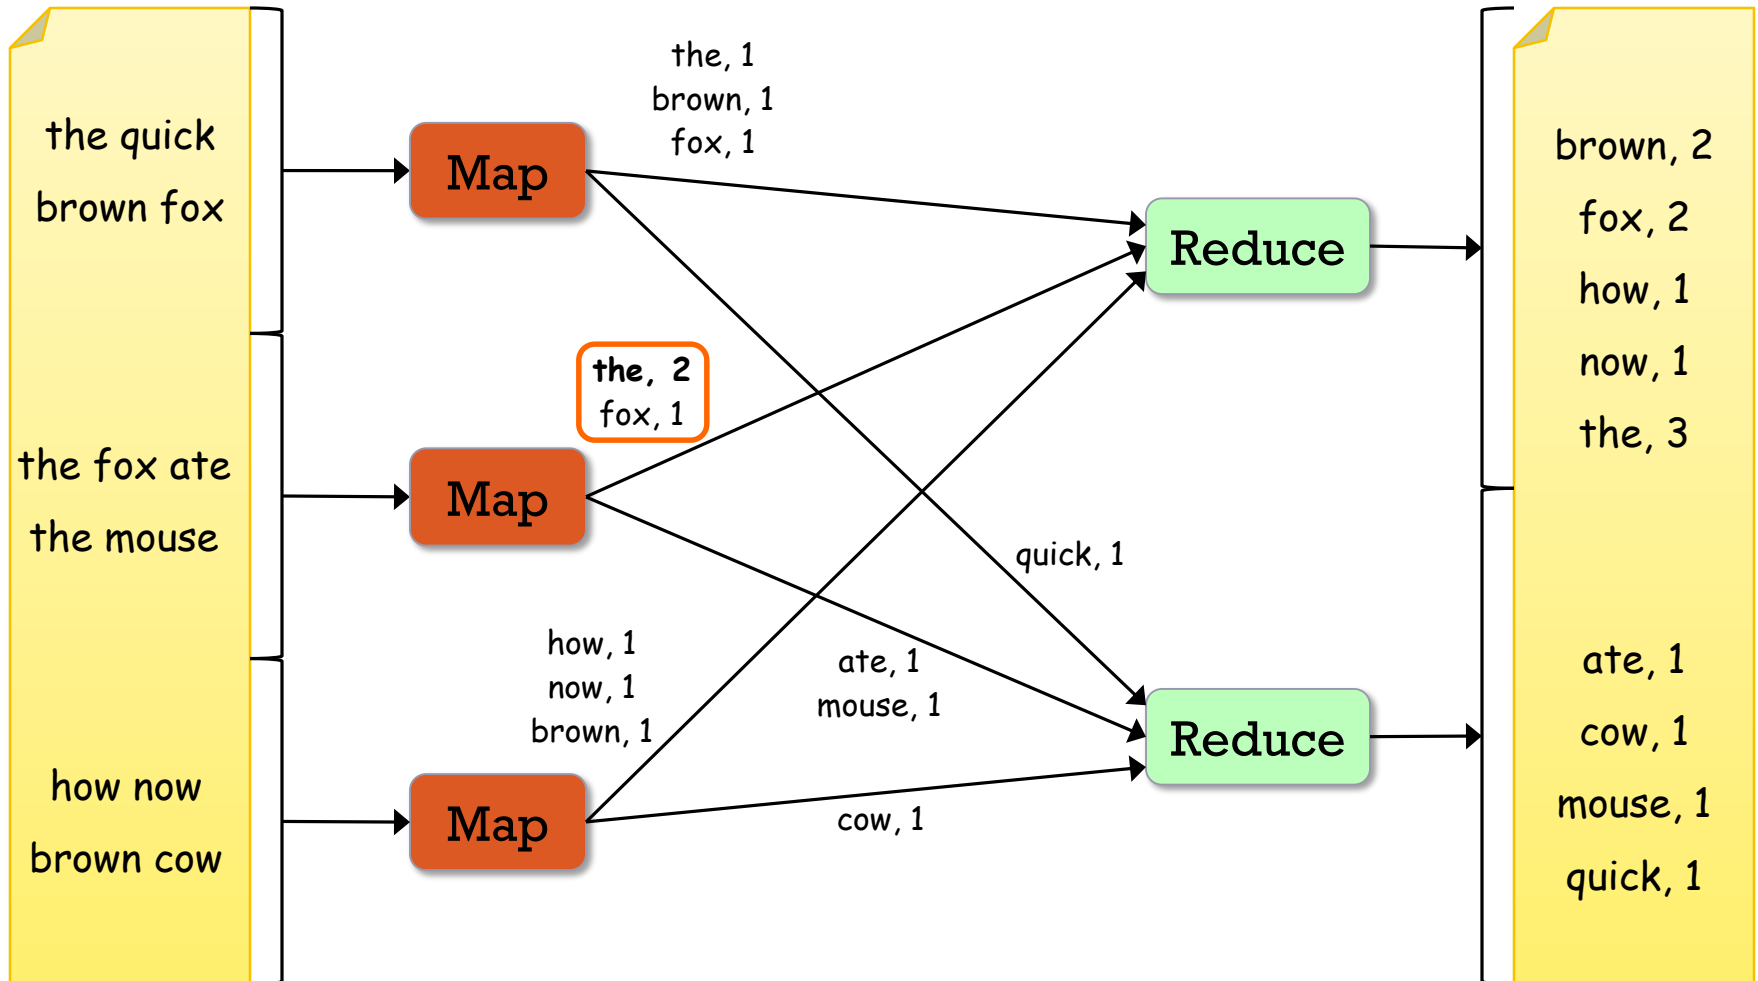
Decreases size of intermediate data

Example: map-side aggregation for Word Count:

```
def combiner(key, values):  
    output(key, sum(values))
```

WORD COUNT WITH COMBINER

Input Map & Combine Shuffle & Sort Reduce Output



FAULT TOLERANCE IN MAPREDUCE

1. If a task crashes:

- Retry on another node
 - OK for a map because it has no dependencies
 - OK for reduce because map outputs are on disk
- If the same task fails repeatedly, fail the job or ignore that input block (user-controlled)

➤ Note: For these fault tolerance features to work, *your map and reduce tasks must be side-effect-free*

FAULT TOLERANCE IN MAPREDUCE

2. If a node crashes:

- Re-launch its current tasks on other nodes
- Re-run any maps the node previously ran
 - Necessary because their output files were lost along with the crashed node

FAULT TOLERANCE IN MAPREDUCE

3. If a task is going slowly (straggler):

- Launch second copy of task on another node (“speculative execution”)
- Take the output of whichever copy finishes first, and kill the other

➤ Surprisingly important in large clusters

- Stragglers occur frequently due to failing hardware, software bugs, misconfiguration, etc
- Single straggler may noticeably slow down a job

TAKEAWAYS

By providing a data-parallel programming model, MapReduce can control job execution in useful ways:

- Automatic division of job into tasks
- Automatic placement of computation near data
- Automatic load balancing
- Recovery from failures & stragglers

User focuses on application, not on complexities of distributed computing

MAPREDUCE PROS

Distribution is completely **transparent**

- Not a single line of distributed programming (ease, correctness)

Automatic **fault-tolerance**

- Determinism enables running failed tasks somewhere else again
- Saved intermediate data enables just re-running failed reducers

Automatic **scaling**

- As operations as side-effect free, they can be distributed to any number of machines dynamically

Automatic **load-balancing**

- Move tasks and speculatively execute duplicate copies of slow tasks (*stragglers*)