

HADOOP

AND OTHERS

INTRODUCTION TO DATA SCIENCE

TIM KRASKA

LAST LECTURE

- **Cloud Computing**
- **HDFS**
- **MapReduce**

THIS LECTURE

- **MapReduce ctd**
- **Other large scale processing frameworks**
- **Small scale processing frameworks**
- **(NO SQL)**

CLICKER

Input to the _____ is the sorted output of the mappers.

- a) Reducer
- b) Mapper
- c) Shuffle
- d) All of the above

MAPREDUCE PROGRAMMING MODEL

Data type: key-value *records*

Map function:

$$(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$$

Reduce function:

$$(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$$

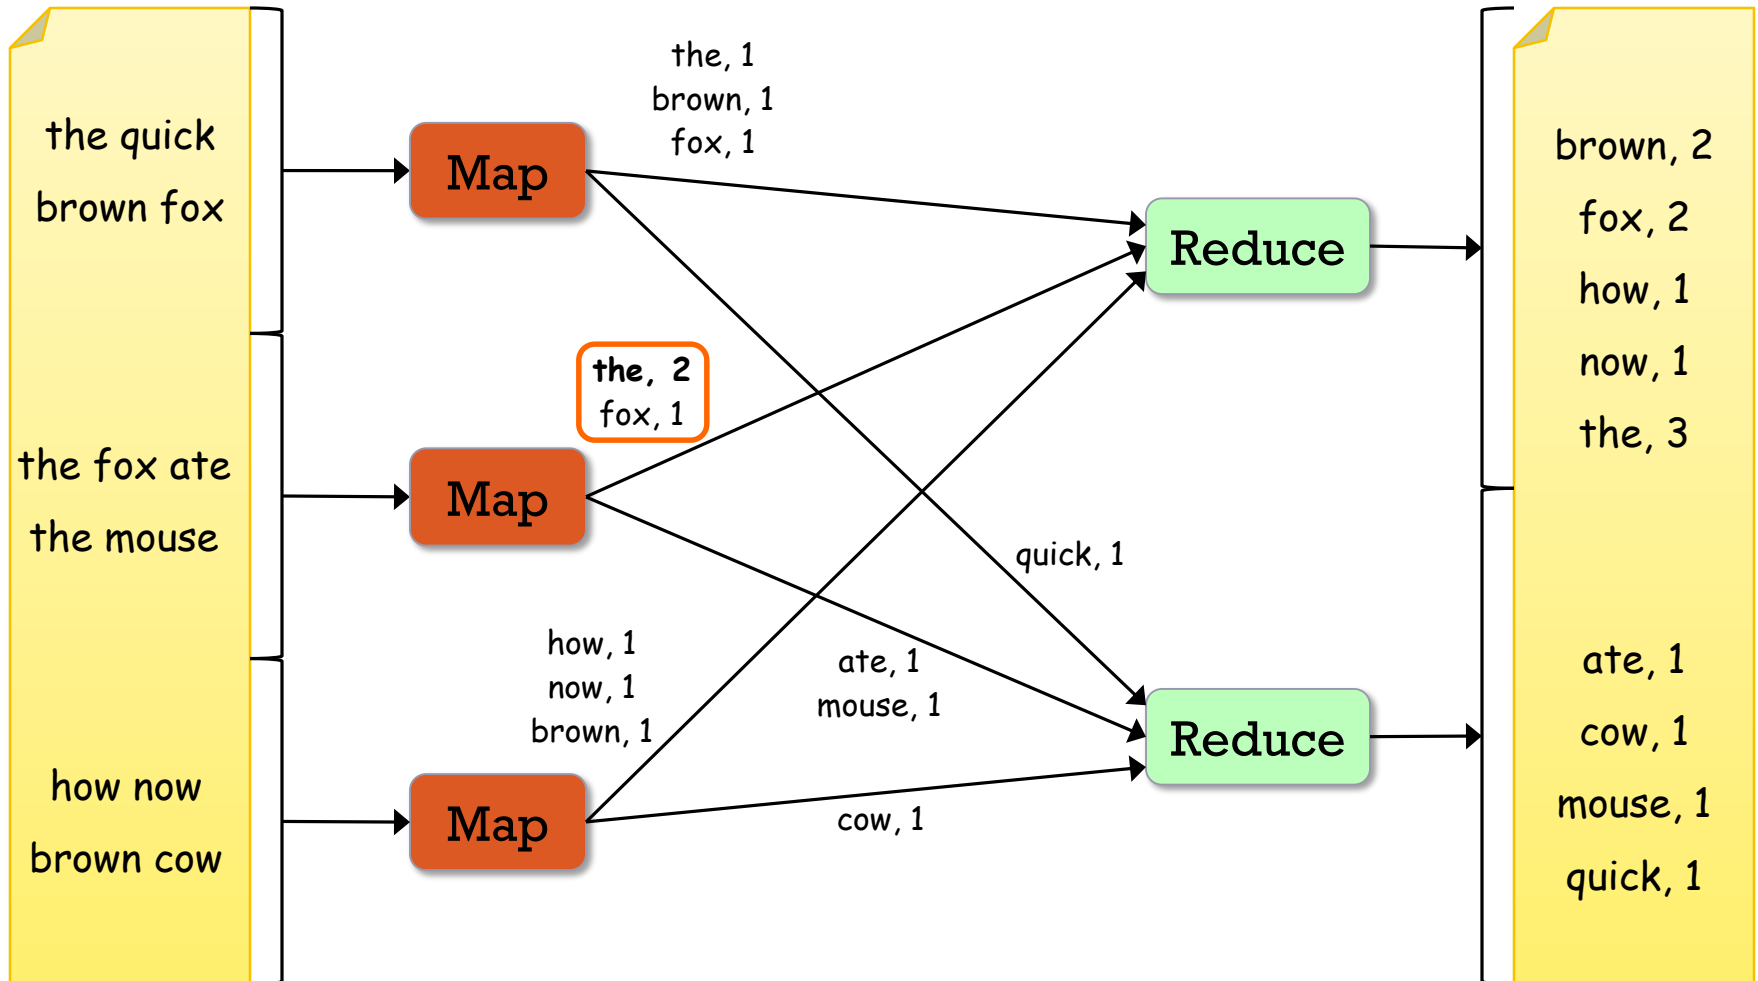
EXAMPLE: WORD COUNT

```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reducer(key, values):  
    output(key, sum(values))
```

WORD COUNT WITH COMBINER


Input Map & Combine Shuffle & Sort Reduce Output








Node 1




Node 2

Node 3









OTHER MAP/REDUCE PARAMETERS

- **One or more Map tasks**
- **Zero or more Reduce tasks**
- **Zero or more Combiner tasks**
- **Shuffle / Partitioning function (distributed)**
- **Sort function (locally executed)**
- **Context for Map,Reduce Combiner**
- **Others (e.g., InputSplit)**
- **Configuration (more on that later)**

MAP/REDUCE PROS

Distribution is completely **transparent**

- Not a single line of distributed programming (ease, correctness)

Automatic **fault-tolerance**

- Determinism enables running failed tasks somewhere else again
- Saved intermediate data enables just re-running failed reducers

Automatic **scaling**

- As operations as side-effect free, they can be distributed to any number of machines dynamically

Automatic **load-balancing**

- Move tasks and speculatively execute duplicate copies of slow tasks (*stragglers*)

A FEW EXAMPLES

1. SEARCH

Input: (lineNumber, line) records

Output: lines matching a given pattern

Map: **if(line matches pattern):**
 output(line)

Reduce: identity function

- Alternative: no reducer (map-only job)

2. SORT

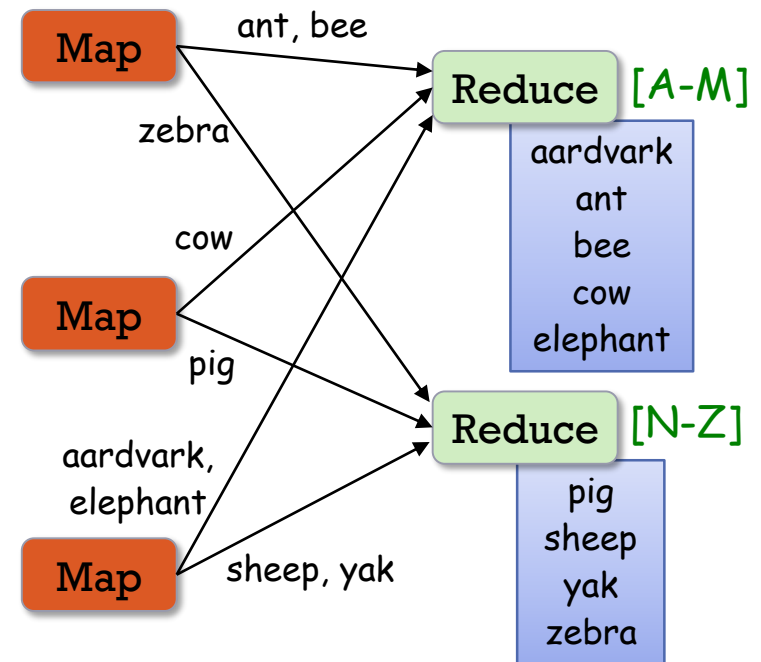
Input: (key, value) records

Output: same records, sorted by key

Map: identity function

Reduce: identity function

Trick: Pick partitioning
function h such that
 $k_1 < k_2 \Rightarrow h(k_1) < h(k_2)$

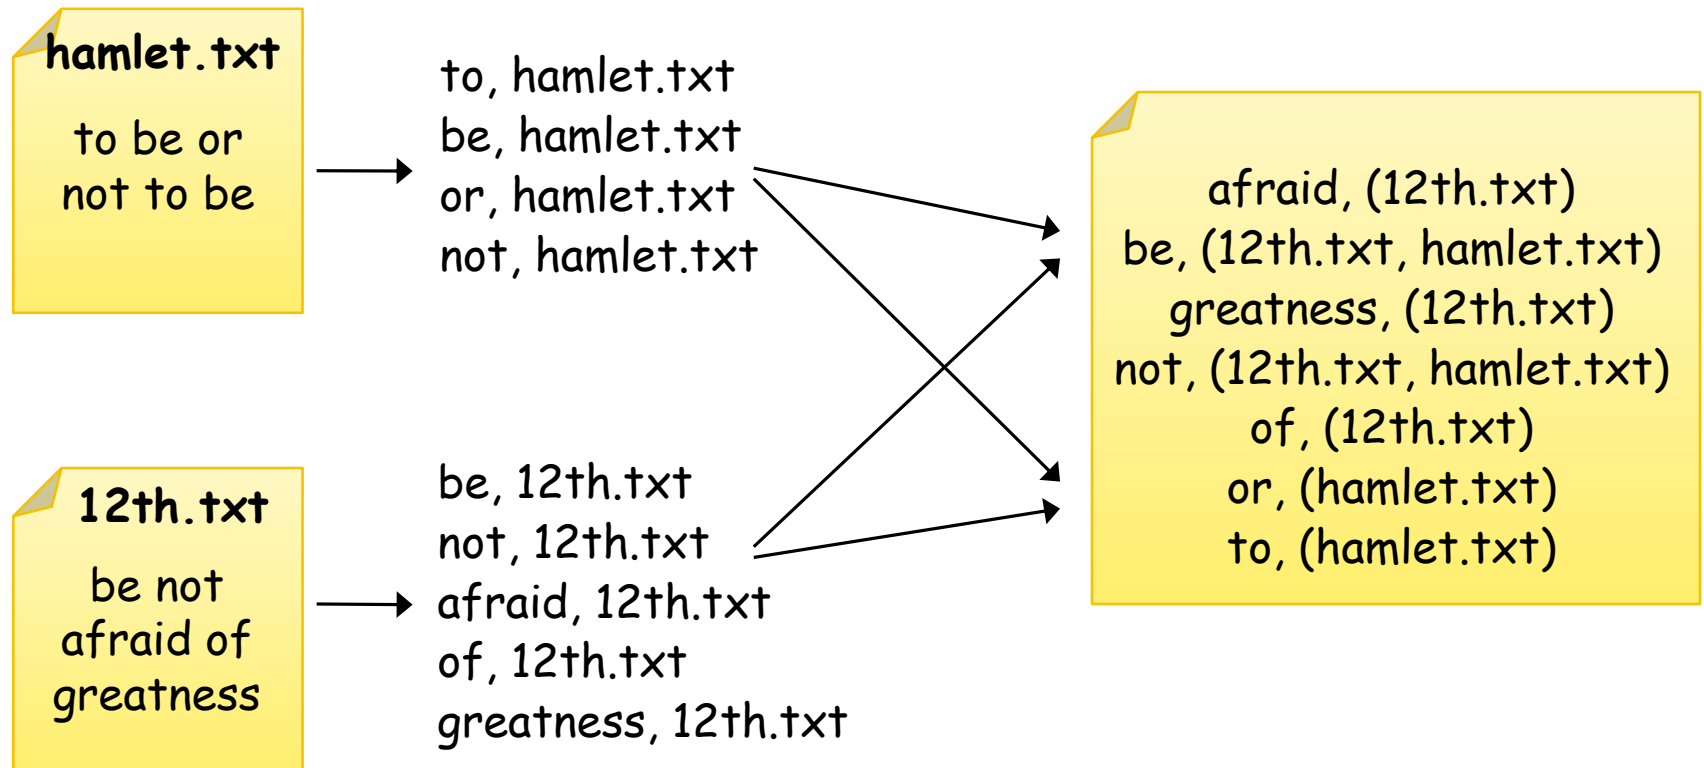


CLICKER: INVERTED INDEX

What MapReduce tasks do you need to build an inverted index

- A) **def Map (filename, text):**
 foreach word in text.split(){
 output(word, filename)}
- def Reduce(word, list(filename)):**
 output(word, sort(filenamees))
- B) **def Map (filename, text):**
 foreach word in text.split(){
 output(word, filename)}
- def Combine(word, filenames):**
 output(word, set(filenames))
- def Reduce(word, filenames):**
 output(word, sort(filenames))
- C) **var globalHashMap = new HashMap on master-node**
def Map (filename, text):
 foreach word in text.split(){
 output(word, filename)}
- def Reduce(word, filenames):**
 globalHashMap.add(word, sort(filenames))

INVERTED INDEX EXAMPLE



3. MOST POPULAR WORDS

Input: (filename, text) records

Output: top 100 words occurring in the most files drop rare words


Two-stage solution:

- **Job 1:**
 - Create inverted index, giving (word, list(file)) records
 - Important: do not remove duplicates
- **Job 2:**
 - Map each (word, list(file)) to (count, word)
 - Sort these records by count as in sort job


Optimizations:

- Map to (word, 1) instead of (word, file) in Job 1
- Count files in job 1's reducer rather than job 2's mapper
- Estimate count distribution in advance and drop rare words

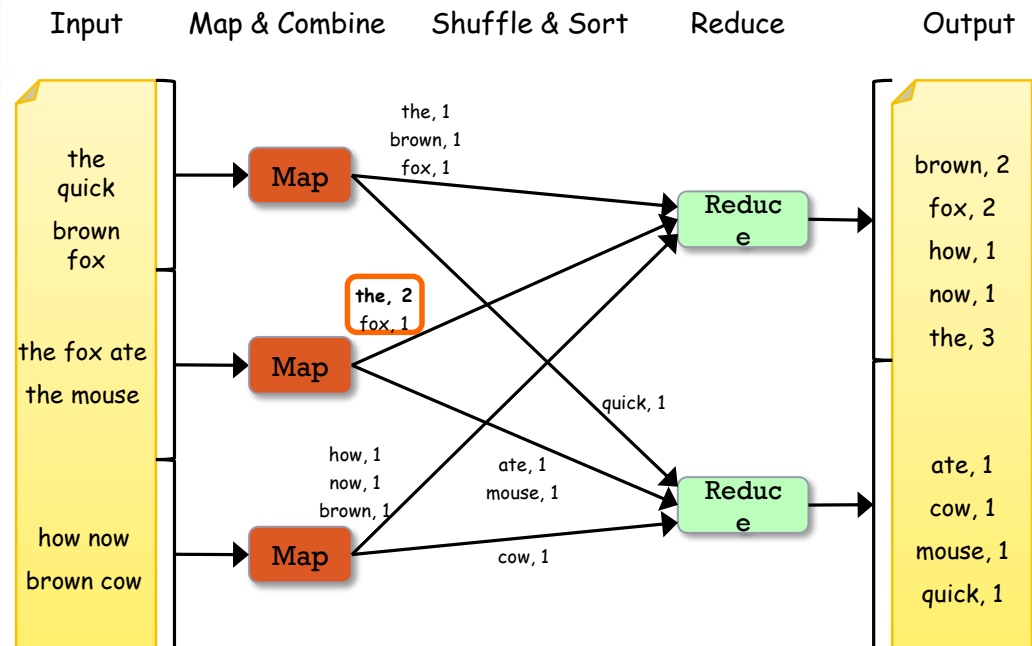

HADOOP: THE 1ST OPEN-SOURCE SYSTEM IMPLEMENTING THE MAPREDUCE PARADIGM



HADOOP ARCHITECTURE



HADOOP ARCHITECTURE



CLICKER

For a very simple word count application on a cluster with 1000 nodes, each having two CPUs, 10 cores each, how many parallel **MAP** tasks (i.e., threads) per node should you use?

- a) 20
- b) $2 * 20$
- c) 60 – 100
- d) More than 100

CLICKER

For a very simple word count application on a cluster with 1000 nodes, each having two CPUs, 10 cores each, how many parallel **REDUCE tasks** per node should you use?

- a) The same number as map tasks
- b) $2 * 1.75 * 20$ per node
- c) $2 * 0.95 * 20$ per node
- d) Needs to be fine tuned so that the output is a multiple of a block size
- e) Needs to be fine tuned so that a reduce task takes between 5 and 10 minutes

Iterative Algorithms in MapReduce

Example KMeans

```
Select  $K$  random data points  $\{s_1, s_2, \dots, s_K\}$  as centroids  $c_j$ .
Until clustering converges or other stopping criterion{
  For each data point  $x_i$ :
    Assign  $x_i$  to the closes centroid such that
       $dist(x_i, c_j)$  is minimal.
  For each cluster  $c_j$ , update the centroids
     $c_j = \mu(c_j)$ 
}
```

How do you express K-Means in the Map/Reduce paradigm?

Iterative Algorithms in MapReduce

Example KMeans

Map1(filename, data) := emit data as (r-id, features)

centroids[] = read-centroids from disk

Configure map2 job with centroids[]

Map2(r-id, features) :=

 compare features (i.e., coordinates) with centroids
 return (Closest-Centroid-ID, features)


Reduce(Centroid-ID, List[features]) :=

 average features (i.e., coordinates) and emit (Centroid-ID, New-Coordinates)


Write new centroids to disk

Check if converged, if not do **Map2** and **Reduce** again

WHAT DO YOU THINK WERE THE REACTION OF THE DATABASE COMMUNITY?



I JUST



**DON'T
CARE**





MR VS. DATABASES

HADOOP VS. RDBMS

Comparison of 3 systems

- Hadoop
- Vertica (a column-oriented database)
- DBMS-X (a row-oriented database)
 - rhymes with “schmoracle”

Qualitative

- Programming model, ease of setup, features, etc.

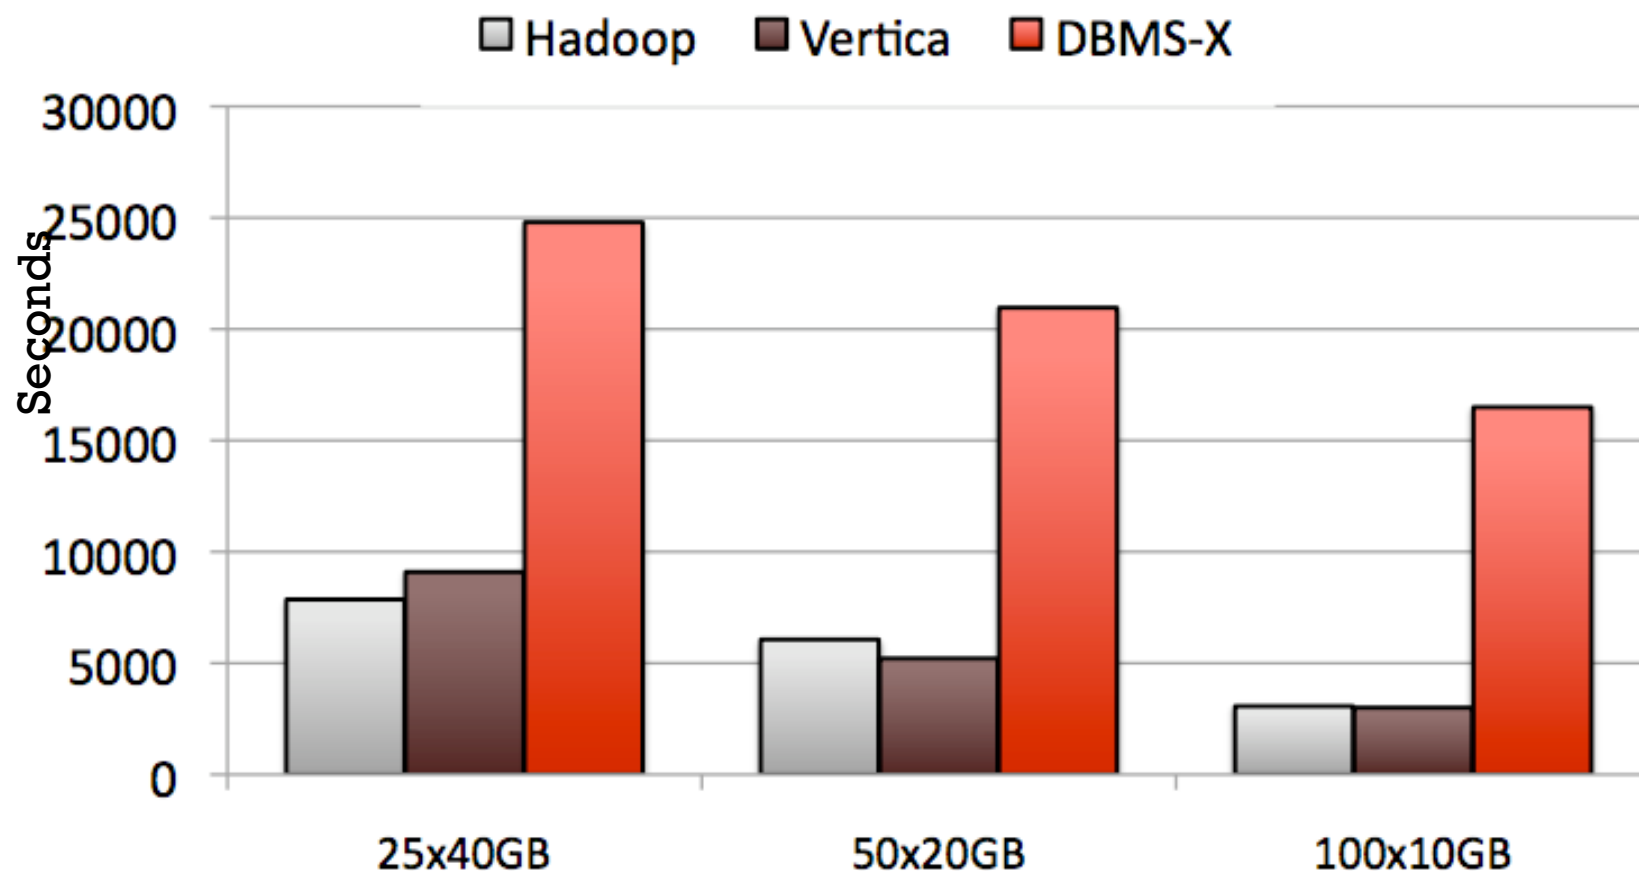
Quantitative

- Data loading, different types of queries


Grep Task

- Find 3-byte pattern in 100-byte record
 - *1 match per 10,000 records*
- Data set:
 - *10-byte unique key, 90-byte value*
 - *1TB spread across 25, 50, or 100 nodes*
 - *10 billion records*
- Original MR Paper (Dean et al. 2004)

Grep Task Loading Results




Grep Task Execution Results



SELECTION TASK

```
SELECT pageURL, pageRank  
FROM Rankings WHERE pageRank > X
```



1 GB /
node

Analytical Tasks


- Simple web processing schema
- Data set:
 - *600k HTML Documents (6GB/node)*
 - *155 million UserVisit records (20GB/node)*
 - *18 million Rankings records (1GB/node)*

Aggregate Task

- Simple query to find adRevenue by IP prefix

```
SELECT SUBSTR(sourceIP, 1, 7),  
       SUM(adRevenue)  
FROM   userVistits  
GROUP BY SUBSTR(sourceIP, 1, 7)
```


Aggregate Task Results



Join Task

- Find the sourceIP that generated the most adRevenue along with its average pageRank.
- Implementations:
 - *DBMSs – Complex SQL using temporary table.*
 - *MapReduce – Three separate MR programs.*

Join Task Results



PROBLEMS WITH THIS ANALYSIS?

Other ways to avoid sequential scans?

Fault-tolerance in large clusters?

Tasks that cannot be expressed as queries?

Google's Response: Cluster Size

- Largest known database installations:
 - *Greenplum – 96 nodes – 4.5 PB (eBay) [1]*
 - *Teradata – 72 nodes – 2+ PB (eBay) [1]*
- Largest known MR installations:
 - *Hadoop – 3658 nodes – 1 PB (Yahoo) [2]*
 - *Hive – 600+ nodes – 2.5 PB (Facebook) [3]*

[1] eBay's two enormous data warehouses – April 30th, 2009

<http://www.dbms2.com/2009/04/30/eBays-two-enormous-data-warehouses/>

[2] Hadoop Sorts a Petabyte in 16.25 Hours and a Terabyte in 62 Seconds – May 11th, 2009

http://developer.yahoo.net/blogs/hadoop/2009/05/hadoop_sorts_a_petabyte_in_162.html

[3] Hive - A Petabyte Scale Data Warehouse using Hadoop – June 10th, 2009

http://www.facebook.com/note.php?note_id=89508453919

Concluding Remarks

- What can *MapReduce* learn from *Databases*?
 - *Declarative languages are a good thing.*
 - *Schemas are important.*
- What can *Databases* learn from *MapReduce*?
 - *Query fault-tolerance.*
 - *Support for in situ data.*
 - *Embrace open-source.*

APACHE PIG

High-level language:

- Expresses sequences of MapReduce jobs
- Provides relational (SQL) operators (JOIN, GROUP BY, etc)
- Easy to plug in Java functions

Started at Yahoo! Research

- Runs about 50% of Yahoo!'s jobs

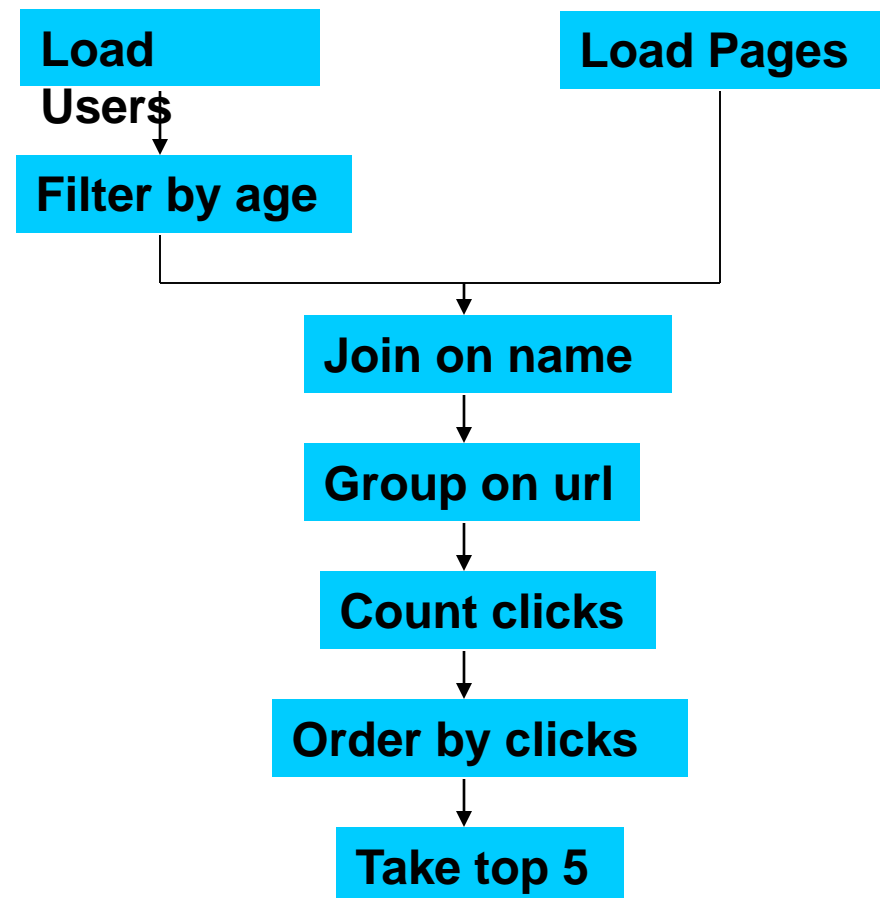
<https://pig.apache.org/>

Similar to Google's (internal) Sawzall project



EXAMPLE PROBLEM

Given *user data* in one file, and *website data* in another, find the *top 5 most visited pages by users aged 18-25*



IN MAPREDUCE

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.Writable;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.RecordReader;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.JobControl.Job;
import org.apache.hadoop.mapred.JobControl.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String key = line.substring(0, firstComma);
            String value = line.substring(firstComma + 1);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("1 " + value);
            oc.collect(outKey, outVal);
        }

        public static class LoadAndFilterUsers extends MapReduceBase
            implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(firstComma + 1);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2 " + value);
            oc.collect(outKey, outVal);
        }

        public static class Join extends MapReduceBase
            implements Reducer<Text, Text, Text, Text> {

        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();

            while (iter.hasNext()) {
                Text t = iter.next();
                String value = t.toString();
                if (value.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }

            reporter.setStatus("OK");
        }

        // Do the cross product and collect the values
        for (String s1 : first) {
            for (String s2 : second) {
                String outval = key + " " + s1 + " " + s2;
                oc.collect(null, new Text(outval));
                reporter.setStatus("OK");
            }
        }
    }

    public static class LoadJoined extends MapReduceBase
        implements Mapper<Text, Text, Text, LongWritable> {

        public void map(
            Text k,
            Text val,
            OutputCollector<Text, LongWritable> oc,
            Reporter reporter) throws IOException {
            // Find the url
            String line = val.toString();
            int firstComma = line.indexOf(',');
            int secondComma = line.indexOf(',', firstComma);
            String key = line.substring(firstComma, secondComma);
            // drop the rest of the record, I don't need it anymore,
            // just pass a 1 for the combiner/reducer to sum instead.
            Text outKey = new Text(key);
            Text outVal = new LongWritable(1L);
            oc.collect(outKey, outVal);
        }

        public static class ReduceUrls extends MapReduceBase
            implements Reducer<Text, LongWritable, WritableComparable,
            Writable> {

        public void reduce(
            Text key,
            Iterator<LongWritable> iter,
            OutputCollector<WritableComparable, Writable> oc,
            Reporter reporter) throws IOException {
            // Add up all the values we see
            long sum = 0;
            while (iter.hasNext()) {
                sum += iter.next().get();
                reporter.setStatus("OK");
            }
            oc.collect(key, new LongWritable(sum));
        }

        public static class LoadClicks extends MapReduceBase
            implements Mapper<WritableComparable, Writable, LongWritable,
            Text> {

        public void map(
            WritableComparable key,
            Writable val,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            oc.collect((LongWritable)val, (Text)key);
        }

        public static class LimitClicks extends MapReduceBase
            implements Reducer<LongWritable, Text, LongWritable, Text> {

        int count = 0;
        public void reduce(
            LongWritable key,
            Iterator<Text> iter,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            // Only output the first 100 records
            while (count < 100 && iter.hasNext()) {
                oc.collect(key, iter.next());
                count++;
            }
        }

        public static void main(String[] args) throws IOException {
            JobConf lp = new JobConf(MRExample.class);
            lp.setJobName("Load Pages");
            lp.setInputFormat(TextInputFormat.class);
            lp.setOutputKeyClass(Text.class);
            lp.setOutputValueClass(Text.class);
            FileInputFormat.addInputPath(lp, new
                Path("/user/gates/pages"));
            FileOutputFormat.setOutputPath(lp, new
                Path("/user/gates/tmp/indexed_pages"));
            lp.setNumReduceTasks(0);
            Job loadPages = new Job(lp);

            JobConf ifu = new JobConf(MRExample.class);
            ifu.setJobName("Load and Filter Users");
            ifu.setInputFormat(TextInputFormat.class);
            ifu.setOutputKeyClass(Text.class);
            ifu.setOutputValueClass(Text.class);
            ifu.setMapperClass(LoadAndFilterUsers.class);
            FileInputFormat.addInputPath(ifu, new
                Path("/user/gates/users"));
            FileOutputFormat.setOutputPath(ifu, new
                Path("/user/gates/tmp/filtered_users"));
            ifu.setNumReduceTasks(0);
            Job loadUsers = new Job(ifu);

            JobConf join = new JobConf(MRExample.class);
            join.setJobName("Join Users and Pages");
            join.setInputFormat(KeyValueTextInputFormat.class);
            join.setOutputKeyClass(Text.class);
            join.setOutputValueClass(Text.class);
            join.setMapperClass(IdentityMapper.class);
            join.setReducerClass(Join.class);
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/indexed_pages"));
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/filtered_users"));
            FileOutputFormat.setOutputPath(join, new
                Path("/user/gates/tmp/joined"));
            join.setNumReduceTasks(50);
            Job joinJob = new Job(join);
            joinJob.addDependingJob(loadPages);
            joinJob.addDependingJob(loadUsers);

            JobConf group = new JobConf(MRExample.class);
            group.setJobName("Group URLs");
            group.setInputFormat(KeyValueTextInputFormat.class);
            group.setOutputKeyClass(Text.class);
            group.setOutputValueClass(LongWritable.class);
            group.setOutputFormat(SequenceFileOutputFormat.class);
            group.setMapperClass(LoadJoined.class);
            group.setCombinerClass(ReduceUrls.class);
            group.setReducerClass(ReduceUrls.class);
            FileInputFormat.addInputPath(group, new
                Path("/user/gates/tmp/joined"));
            FileOutputFormat.setOutputPath(group, new
                Path("/user/gates/tmp/grouped"));
            group.setNumReduceTasks(50);
            Job groupJob = new Job(group);
            groupJob.addDependingJob(joinJob);

            JobConf top100 = new JobConf(MRExample.class);
            top100.setJobName("Top 100 sites");
            top100.setInputFormat(SequenceFileInputFormat.class);
            top100.setOutputKeyClass(LongWritable.class);
            top100.setOutputValueClass(Text.class);
            top100.setOutputFormat(SequenceFileOutputFormat.class);
            top100.setMapperClass(LoadClicks.class);
            top100.setCombinerClass(LimitClicks.class);
            top100.setReducerClass(LimitClicks.class);
            FileInputFormat.addInputPath(top100, new
                Path("/user/gates/tmp/grouped"));
            FileOutputFormat.setOutputPath(top100, new
                Path("/user/gates/top100sitesforusers18to25"));
            top100.setNumReduceTasks(1);
            Job limit = new Job(top100);
            limit.addDependingJob(groupJob);

            JobControl jc = new JobControl("Find top 100 sites for users
                18 to 25");
            jc.addJob(loadPages);
            jc.addJob(loadUsers);
            jc.addJob(joinJob);
            jc.addJob(groupJob);
            jc.addJob(limit);
            jc.run();
        }
    }
}
```

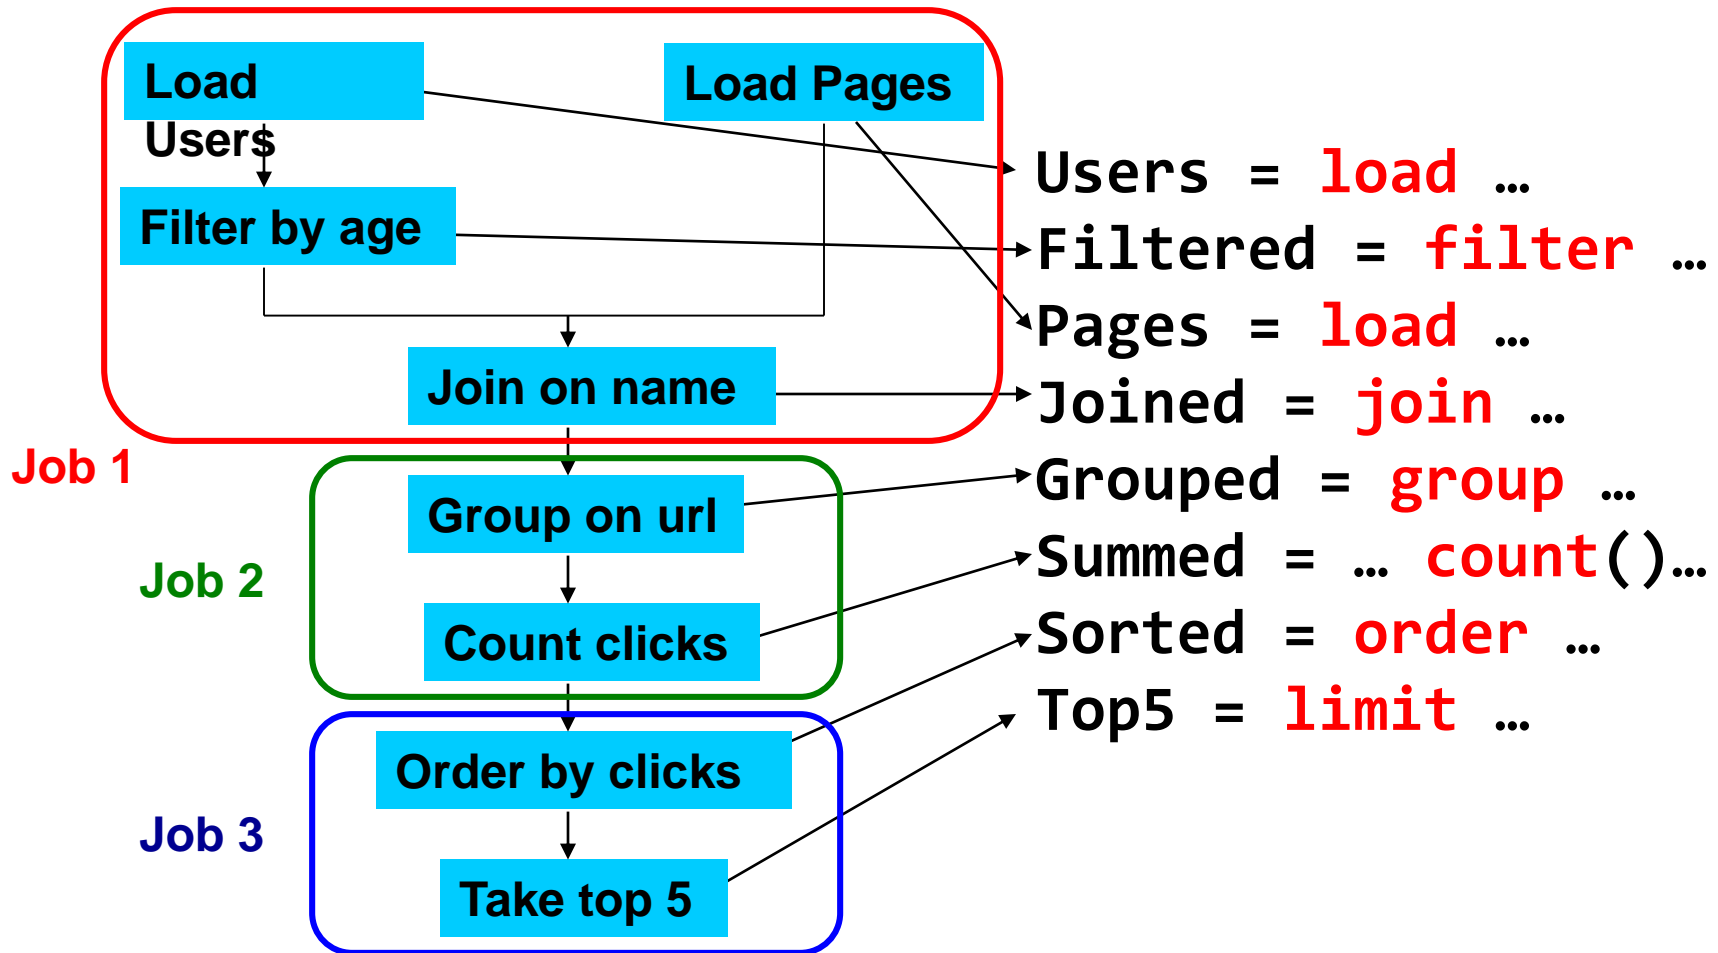
IN PIG LATIN

```
Users      = load 'users' as (name, age);
Filtered   = filter Users by
              age >= 18 and age <= 25;
Pages      = load 'pages' as (user, url);
Joined     = join Filtered by name, Pages by user;
Grouped    = group Joined by url;
Summed     = foreach Grouped generate group,
              count(Joined) as clicks;
Sorted     = order Summed by clicks desc;
Top5       = limit Sorted 5;

store Top5 into 'top5sites';
```


TRANSLATION TO MAPREDUCE

Notice how naturally the components of the job translate into Pig Latin



APACHE HIVE

Relational database built on Hadoop

- Maintains table schemas
- SQL-like query language (which can also call Hadoop Streaming scripts)
- Supports table partitioning, complex data types, sampling, some query optimization

Developed at Facebook

- Used for many Facebook jobs

Now used by many others

- Netflix, Amazon, ...

<http://hive.apache.org/>






APACHE SPARK MOTIVATION


Lightning-Fast Cluster Computing

Complex jobs, interactive queries and online processing all need one thing that MR lacks:


Efficient primitives for data sharing



Iterative job



Interactive mining




Stream processing

SPARK MOTIVATION

Complex jobs, interactive queries and online processing all need one thing that MR lacks:

Efficient primitives for data sharing



The diagram shows a horizontal bar with several gray rectangular blocks representing jobs. Above the bar, there are three gray blocks on the left and two yellow cylindrical blocks on the right. A label 'Query 1' is positioned above the middle of the bar. A yellow rounded rectangle highlights the text below.



Problem: in MR, the only way to share data across jobs is using stable storage (e.g. file system) → slow!

Iterative job

Interactive mining

Stream processing


EXAMPLES



Opportunity: DRAM is getting cheaper →
use main memory for intermediate
results instead of disks

...

GOAL: IN-MEMORY DATA SHARING



10-100 × faster than network and disk

SPARK PROGRAMMING MODEL

Resilient distributed datasets (RDDs)

- Immutable collections partitioned across cluster that can be rebuilt if a partition is lost
- Created by transforming data in stable storage using data flow operators (map, filter, group-by, ...)
- Can be *cached* across parallel operations

Parallel operations on RDDs

- Reduce, collect, count, save, ...

Restricted shared variables

- Accumulators, broadcast variables

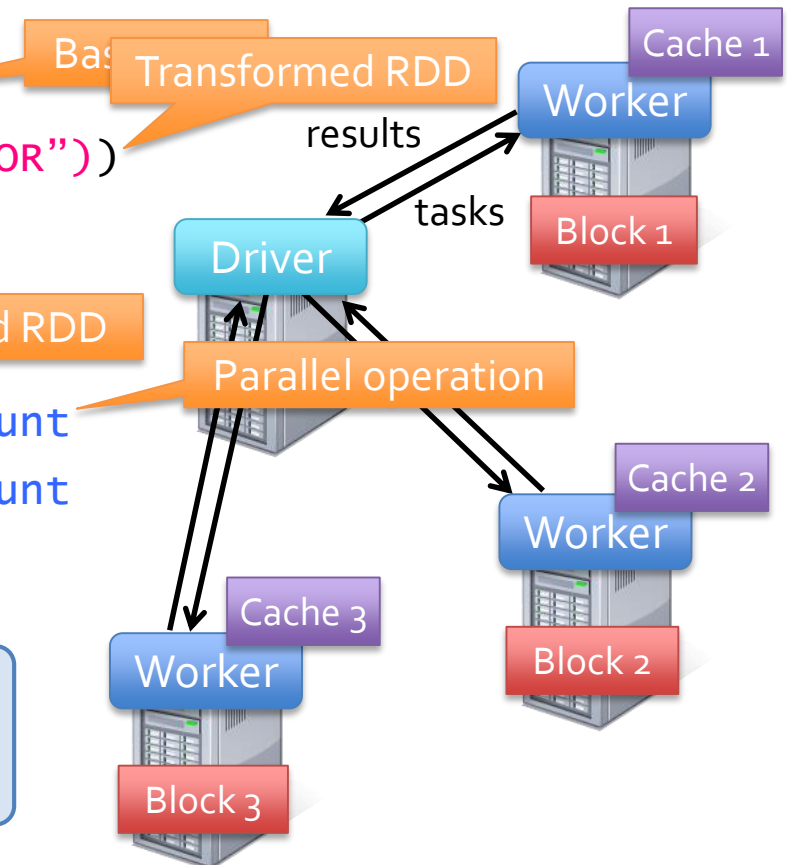
Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()

cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

Result: full-text search of Wikipedia
in <1 sec (vs 20 sec for on-disk data)



RDDS IN MORE DETAIL

An RDD is an immutable, partitioned, logical collection of records

- Need not be materialized, but rather contains information to rebuild a dataset from stable storage

Partitioning can be based on a key in each record (using hash or range partitioning)

Built using bulk transformations on other RDDs

Can be cached for future reuse

RDD OPERATIONS

Transformations (define a new RDD)

map
filter
sample
union
groupByKey
reduceByKey
join
cache
...

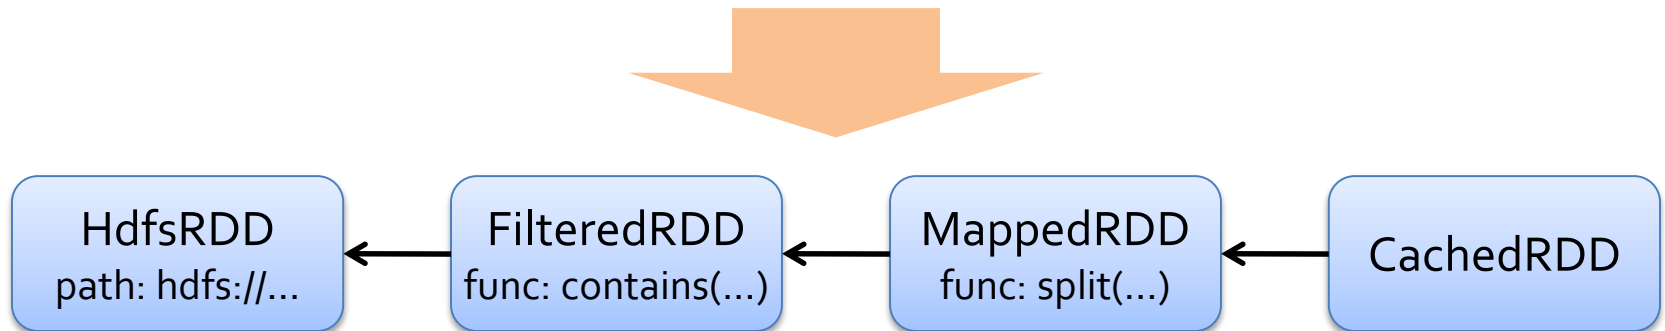
Parallel operations (return a result to driver)

reduce
collect
count
save
lookupKey
...

RDD FAULT TOLERANCE

RDDs maintain *lineage* information that can be used to reconstruct lost partitions

Ex: `cachedMsgs = textFile(...).filter(_.contains("error")).map(_.split('\t')(2)).cache()`



BENEFITS OF RDD MODEL

Consistency is easy due to immutability

Inexpensive fault tolerance (log lineage rather than replicating/checkpointing data)

Locality-aware scheduling of tasks on partitions

Despite being restricted, model seems applicable to a broad variety of applications

RDDS VS DISTRIBUTED SHARED MEMORY

Concern	RDDs	Distr. Shared Mem.
Reads	Fine-grained	Fine-grained
Writes	Bulk transformations	Fine-grained
Consistency	Trivial (immutable)	Up to app / runtime
Fault recovery	Fine-grained and low-overhead using lineage	Requires checkpoints and program rollback
Straggler mitigation	Possible using speculative execution	Difficult
Work placement	Automatic based on data locality	Up to app (but runtime aims for transparency)