

SQL

INTRODUCTION TO DATA SCIENCE

TIM KRASKA



THE 1ST PROJECT

- Goal: get started to develop a data product
- Projects can range from gathering and cleaning a data set, over repeating an existing study or web-site to visualize a data set to testing an entirely new hypothesis.
- Groups of 4 (use Piazza if you still need a group)

2/16/2016	Pre-Proposal Handin
3/3/2016	Advisor Checkin
3/22/2016	Mid-term report
4/12/2016	Full Project Proposal
4/21/2016	Advisor Checkin
5/3/2016	Advisor Checkin
5/12/2016	Final Project Due

- Mid-term report has to be a public blog post
- 2nd project can be a continuations of the 1st, but doesn't have to be
- 10% of your grade (the final project will be 25%).
- **List of potential data sources and some project ideas are online**
http://cs.brown.edu/courses/csci1951-a/final_project.shtml
- **You are encouraged to find your own data and create your own project**
- Pre-proposal is not binding even for the mid-term report
- We will have a competition for the best Mid-term report/project

LAST LECTURE

- SQL basics:

Select
From
Where
Group By
Having
Order By

- Nested queries
- Set operations
- Null values

CLICKER: EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name
from Professor p
where not exists ( select *
                  from Lecture l
                  where l.Person-ID = p.Person-ID );
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

What is the result of the query?

- (a) {Stan, Eli}
- (b) {Ugur}
- (c) {Eli}

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name  
from Professor p  
where not exists ( select *  
                   from Lecture l  
                   where l.Person-ID = p.Person-ID );
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name
from Professor p
where not exists ( select *
                  from Lecture l
                  where l.Person-ID = 1);
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name  
from Professor p  
where not exists ( select *  
                   from Lecture l  
                   where l.Person-ID = p.Person-ID );
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name
from Professor p
where not exists ( select *
                  from Lecture l
                  where l.Person-ID = 2);
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name  
from Professor p  
where not exists ( select *  
                   from Lecture l  
                   where l.Person-ID = 2);
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

Result: {Ugur}

EXISTENTIAL QUANTIFICATION: EXISTS SUB-QUERIES

```
select p.Name
from Professor p
where not exists ( select *
                  from Lecture l
                  where l.Person-ID = 3);
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

Result: {Ugur}

UNCORRELATED SUB-QUERY

```
select Name  
from Professor  
where Person-ID not in ( select Person-ID  
                           from Lecture);
```

Professor

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Lecture

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

UNCORRELATED SUB-QUERY

```
select Name  
from Professor  
where Person-ID not in ( select Person-ID  
                           from Lecture);
```

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

UNCORRELATED SUB-QUERY

```
select Name  
from Professor  
where Person-ID not in ( select Person-ID  
                           from Lecture);
```

Person-Id	Name
1	Stan
2	Ugur
3	Eli

Person-Id	Course-ID
1	CS127
1	CS227
3	CS155

SYNTACTIC SUGAR

```
select *  
from Student  
where 1 <= Semester and Semester < = 6;
```

```
select *  
from Student  
where Semester between 1 and 6;
```

```
select *  
from Student  
where Semester in (2,4,6);
```

COMPARISONS WITH LIKE

"%,, represents any sequence of characters (0 to n)

"_,, represents exactly one character

N.B.: For comparisons with = , % and _ are normal chars.

```
select *
```

```
from Student
```

```
where Name like 'Tim%';
```

```
select distinct Name
```

```
from Lecture l, attends a, Student s
```

```
where s.Student-ID = a.Student-ID and a.CID = l.CID  
      and l.Title like '%science%';
```

CLICKER

Given the following tables

Runners

id	name
1	John
2	Tim
3	Alice
4	Lisa

Races

Event_id	Event	Winner_id
1	Tough mudder	2
2	500m	3
3	Cross-country	2
4	Triathlon	null

What will be the result of the query:

```
SELECT *
```

```
FROM runners
```

```
WHERE id NOT IN (SELECT winner_id FROM races)
```

- A) 1
- B) Empty set
- C) (1,4)

JOINS IN SQL-92

- **cross join**: Cartesian product
- **natural join**
- **join** or **inner join**
- **left, right** or **full outer join**
- (union join: not discussed here)

```
select *  
from R1, R2  
where R1.A = R2.B;
```

```
select *  
from R1 join R2 on R1.A = R2.B;
```

LEFT OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID,  
s.Student-ID, s.Name  
from Professor p left outer join  
    (tests t left outer join Student s  
    on t.Student-ID= s.Student-ID)  
    on p.Person-ID=t.Person-ID;
```

Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
2136	Curie	-	-	-	-	-
:	:	:	:	:	:	:

RIGHT OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID, s.Student-ID, s.Name
from Professor p right outer join
    (tests t right outer join Student s on
        t.Student-ID= s.Student-ID)
on p.Person-ID=t.Person-ID;
```

Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
-	-	-	-	-	26120	Fichte
:	:	:	:	:	:	:

FULL OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID,  
s.Student-ID, s.Name
```

```
from Professor p full outer join
```

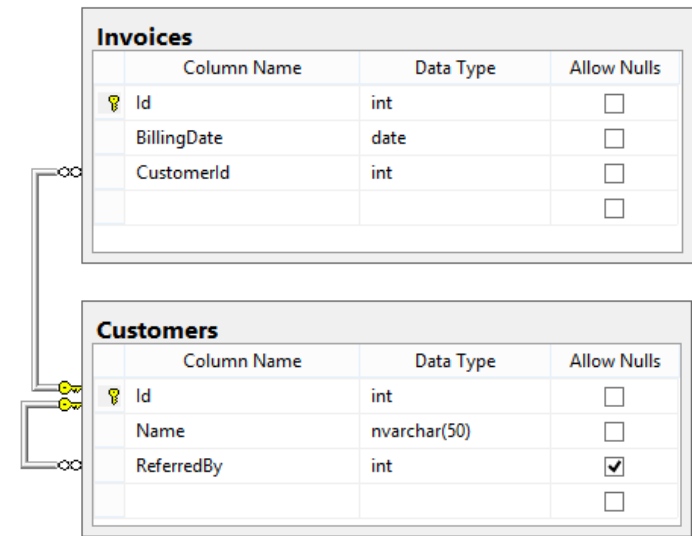
```
    (tests t full outer join Student s on  
    t.Student-ID= s.Student-ID)
```

```
on p.Person-ID=t.Person-ID;
```

p.Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
-	-	-	-	-	26120	Fichte
2136	Curie	-	-	-	-	-

CLICKER

You have your first day as a Data Scientist at TheShop LLC. TheShop uses SQL server. A quick look using VisualStudio shows the simply database schema on the right. In order to analyze all the orders in Python, your first task is to write a SQL query to return a list of all the invoices. For each invoice, you want the Invoice ID, the billing date, the customer's name, and the name of the customer who referred that customer (if any). The list should be ordered by billing date.



- A) `SELECT i.Id, i.BillingDate, c.Name, r.Name AS ReferredByName
FROM (Invoices i JOIN Customers c ON i.CustomerId = c.Id)
JOIN Customers r ON c.ReferredBy = r.Id
ORDER BY i.BillingDate;`
- B) `SELECT i.Id, i.BillingDate, c.Name, r.Name AS ReferredByName
FROM (Invoices I JOIN Customers c ON i.CustomerId = c.Id)
LEFT JOIN Customers r ON c.ReferredBy = r.Id
ORDER BY i.BillingDate;`
- C) `SELECT i.Id, i.BillingDate, c.Name, r.Name AS ReferredByName
FROM (Invoices i Right JOIN Customers c ON i.CustomerId = c.Id)
Right JOIN Customers r ON c.ReferredBy = r.Id
ORDER BY i.BillingDate;`

HOW DOES THE DB PROCESS A SQL QUERY?

SQL is the “WHAT” not the “HOW”

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)


```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = y.cid and
      x.price > 100 and z.city = 'Providence'
```

It's clear WHAT we want, unclear HOW to get it

Exposing the Algebra: Microsoft SQL Server

```
1 select a.term_id, b.term_id
2 from [billhowe].[reuters] a, [billhowe].[reuters] b
3 where a.doc_id = b.doc_id
4      and a.term_id != b.term_id
```


EXPLAIN



Exposing the Algebra: Microsoft SQL Server


```
1 select a.term_id, b.term_id
2 from [billhowe].[reuters] a, [billhowe].[reuters] b
3 where a.doc_id = b.doc_id
4       and a.term_id != b.term_id
5       and a.term_id = 'parliament'
```

EXPLAIN



Exposing the Algebra: PostgreSQL

EXPLAIN SELECT



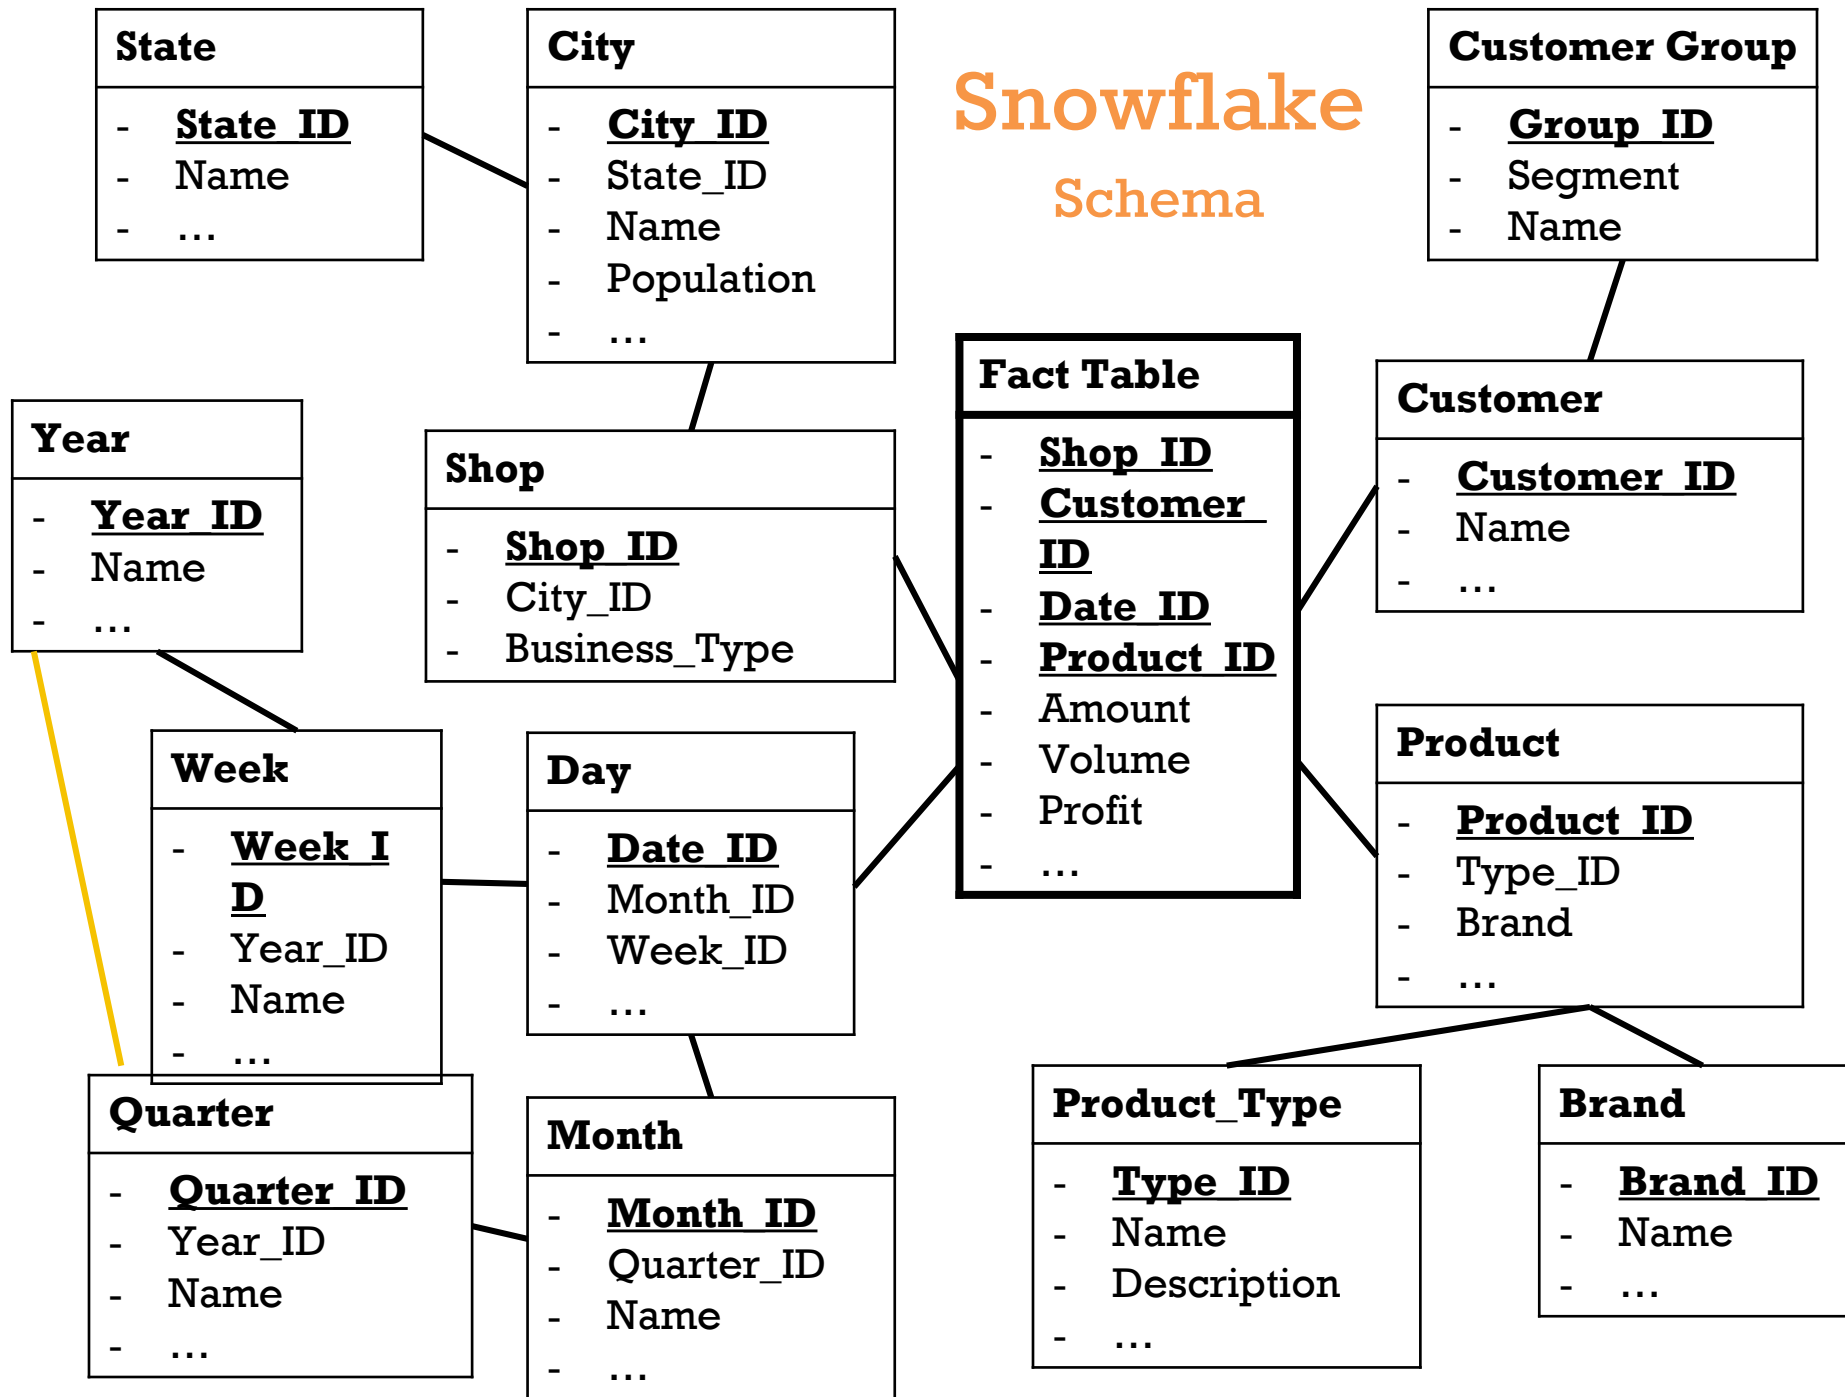
screenshot from pgAdmin3

STAR VS SNOWFLAKE SCHEMA

DATA WAREHOUSES

Snowflake

Schema



Star

Schem

a

Shop
- <u>Shop ID</u>
- Business_Type
- City
- City_Population
- State
- ...

Customer
- <u>Customer ID</u>
- Name
- Segment
- Group_Name
- ...

Fact Table
- <u>Shop ID</u>
- <u>Customer ID</u>
- <u>Date ID</u>
- <u>Product ID</u>
- Amount
- Volume
- Profit
- ...

Time
- <u>Date ID</u>
- Month
- Quarter
- Year
- ...

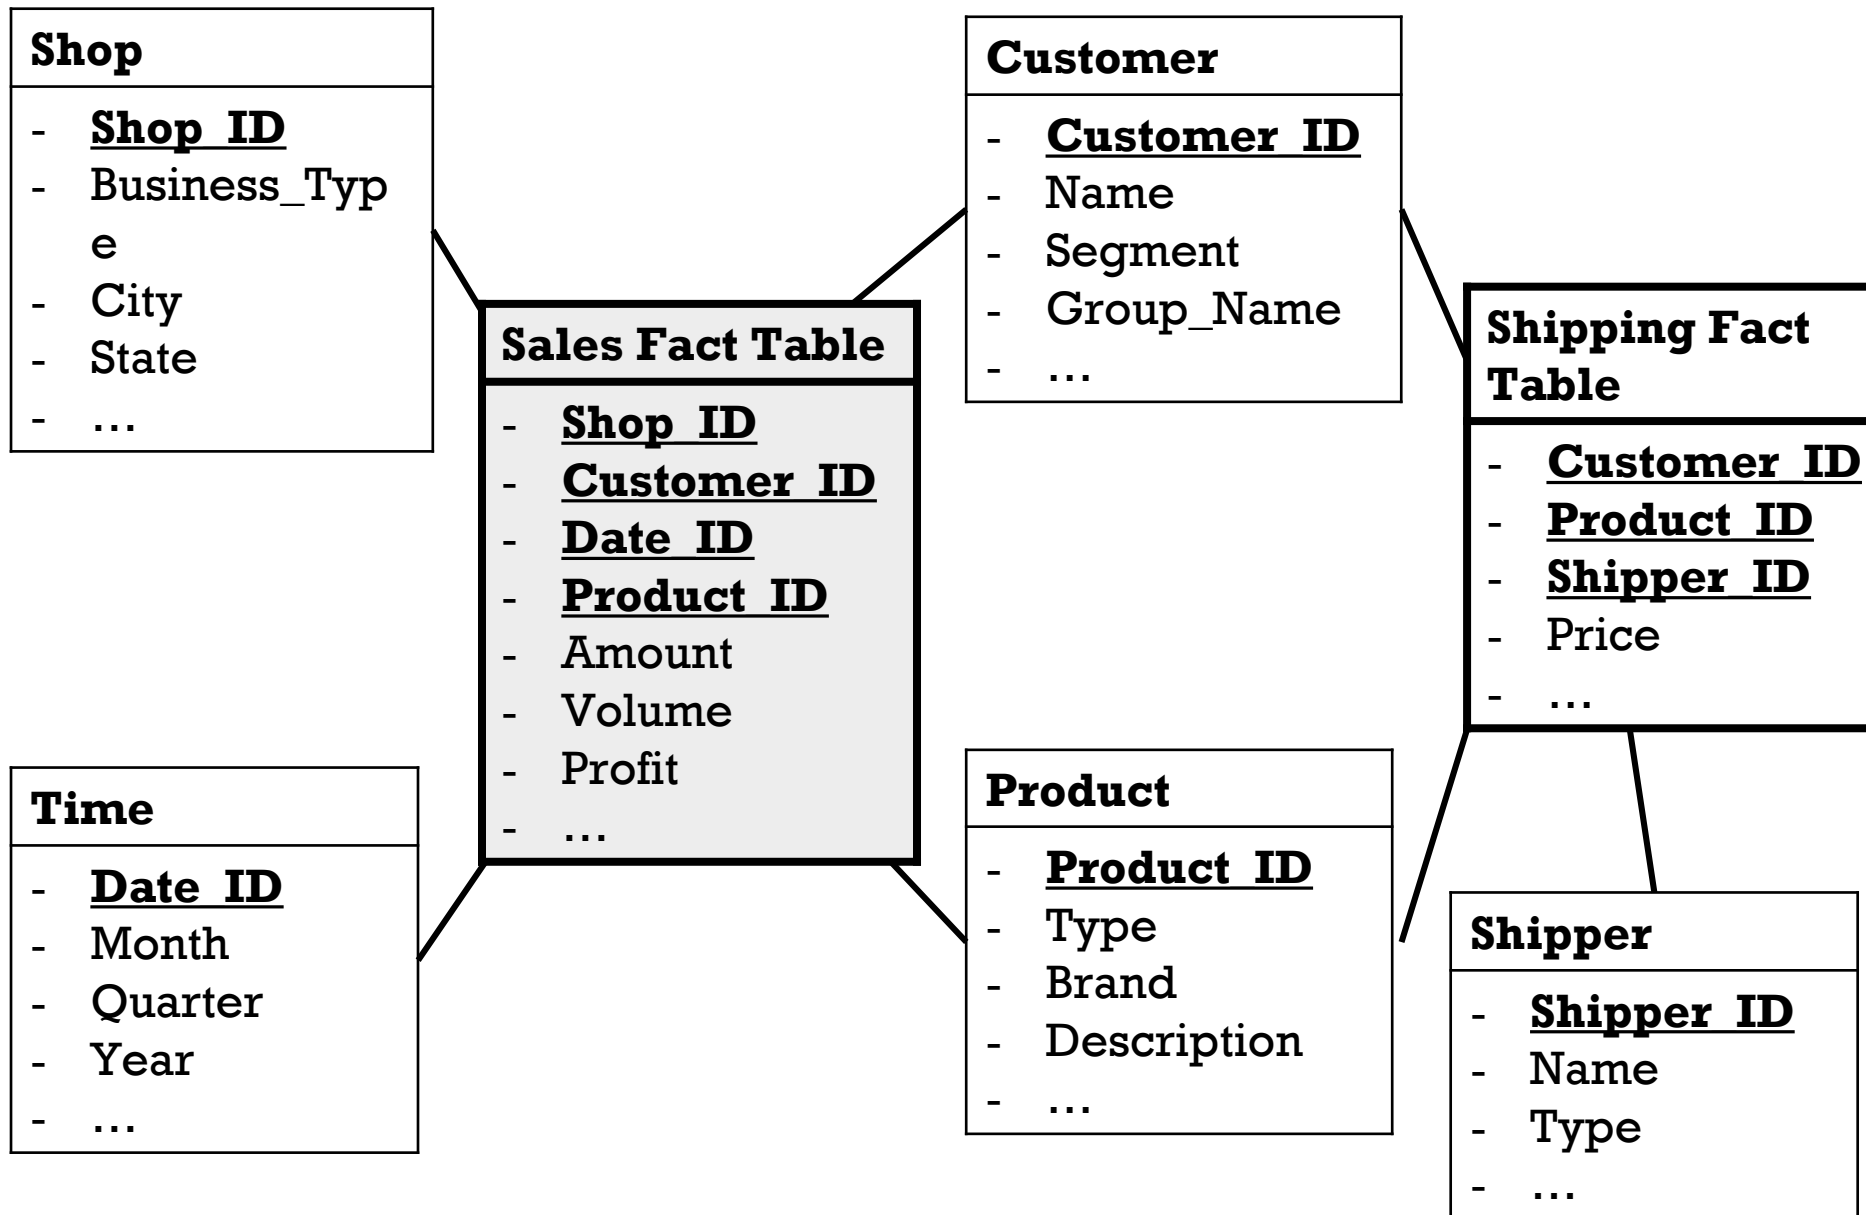
Product
- <u>Product ID</u>
- Type
- Brand
- Description
- ...

STAR VS. SNOWFLAKE SCHEMA

	Snowflake	Star
Normalization/ De-Normalization	Dimension Tables are in Normalized form but Fact Table is still in De-Normalized form	Both Dimension and Fact Tables are in De-Normalized form
Space	Smaller	Bigger (Redundancy)
Query Performance	More Joins → slower	Fewer Joins → faster
Ease of Use	Complex Queries	Pretty Simply Queries
When to use	When dimension table is relatively big in size, snowflaking is better as it reduces space.	When dimension table contains less number of rows, we can go for Star schema.

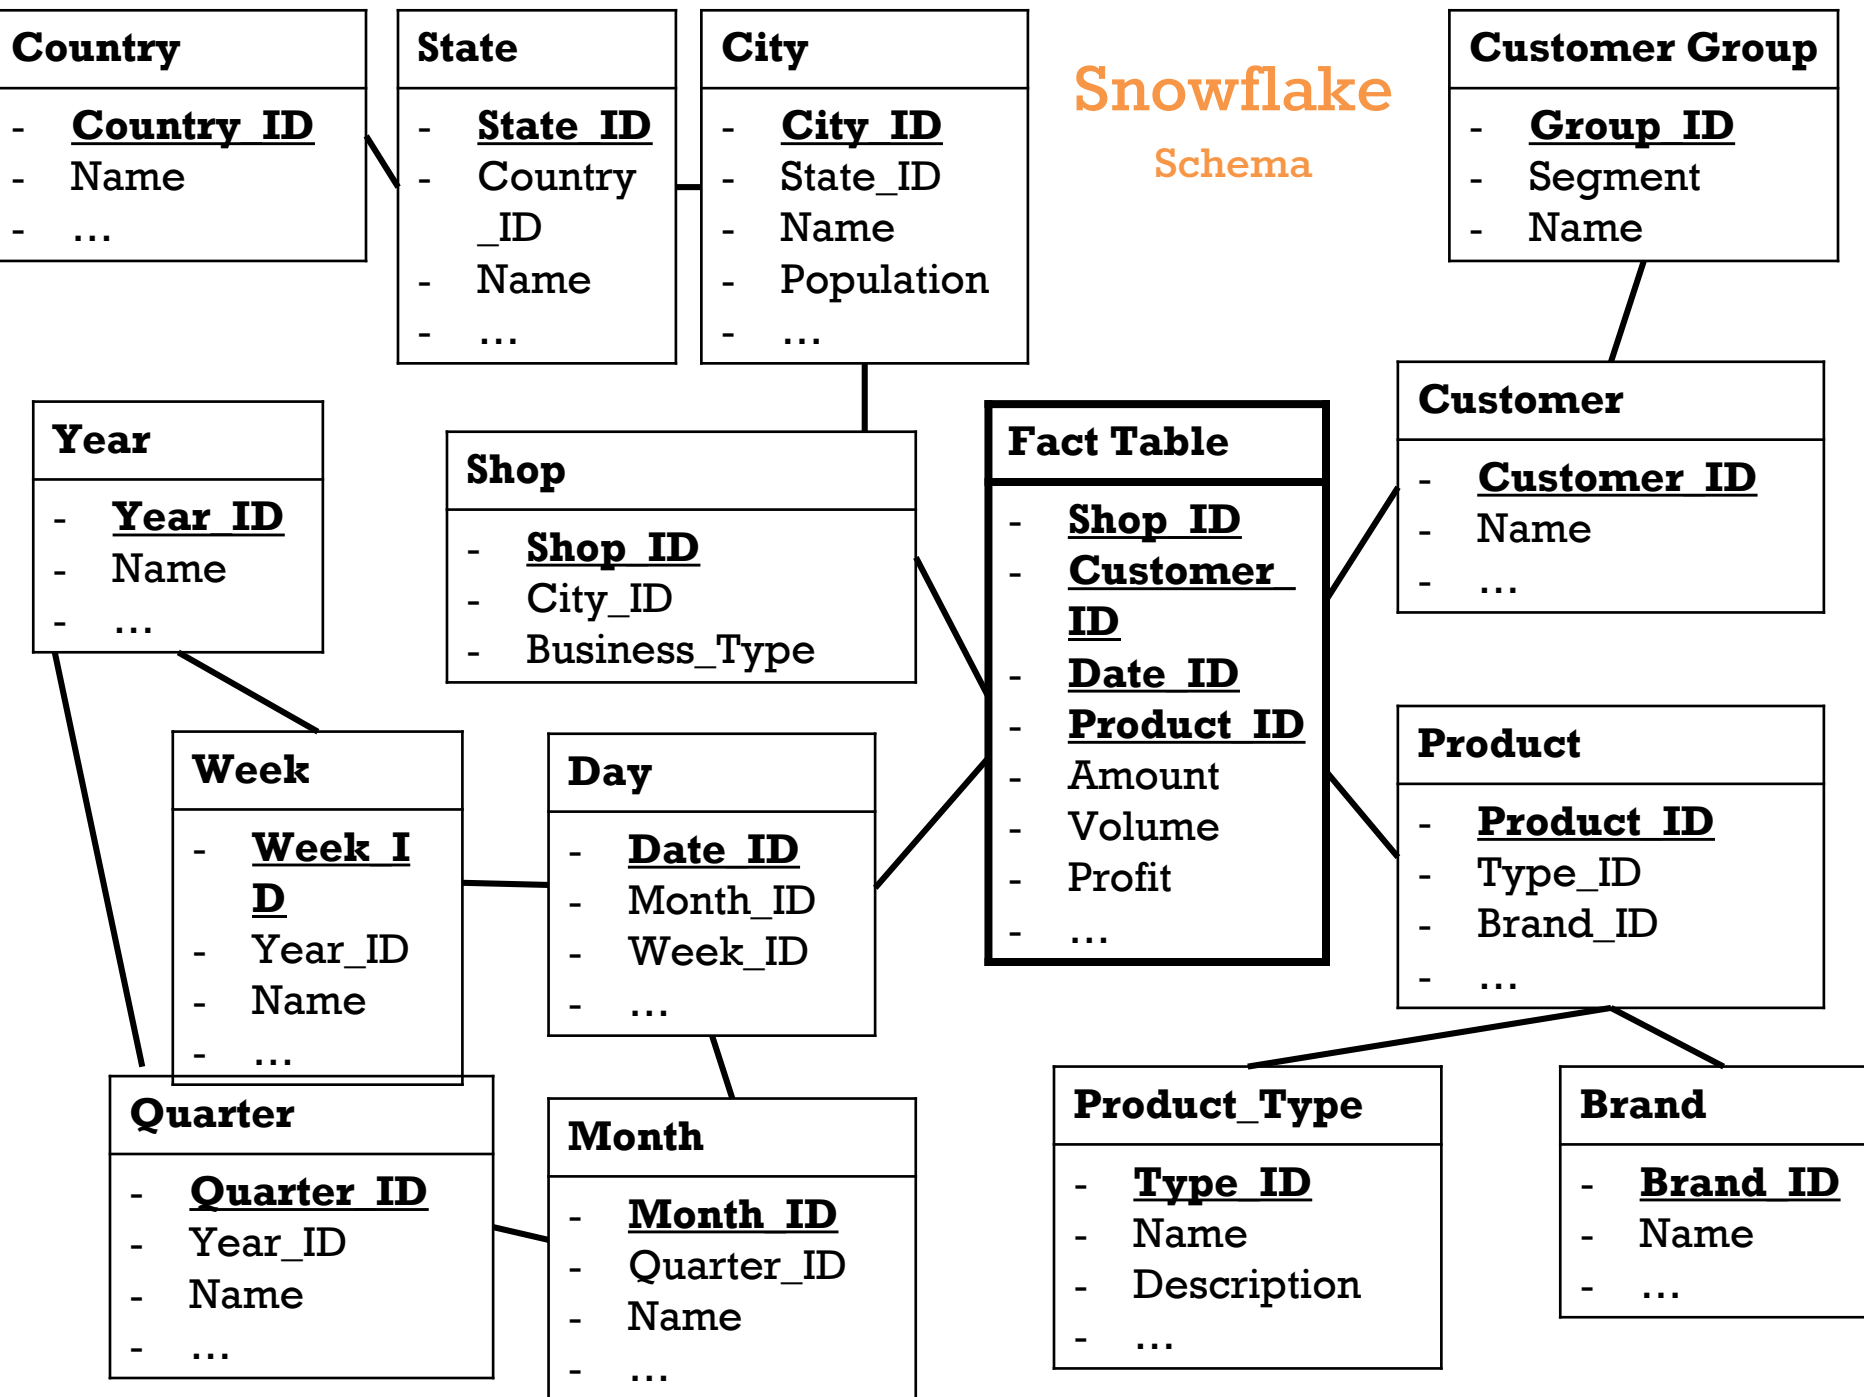
Galaxy / Fact Constellation

Schema

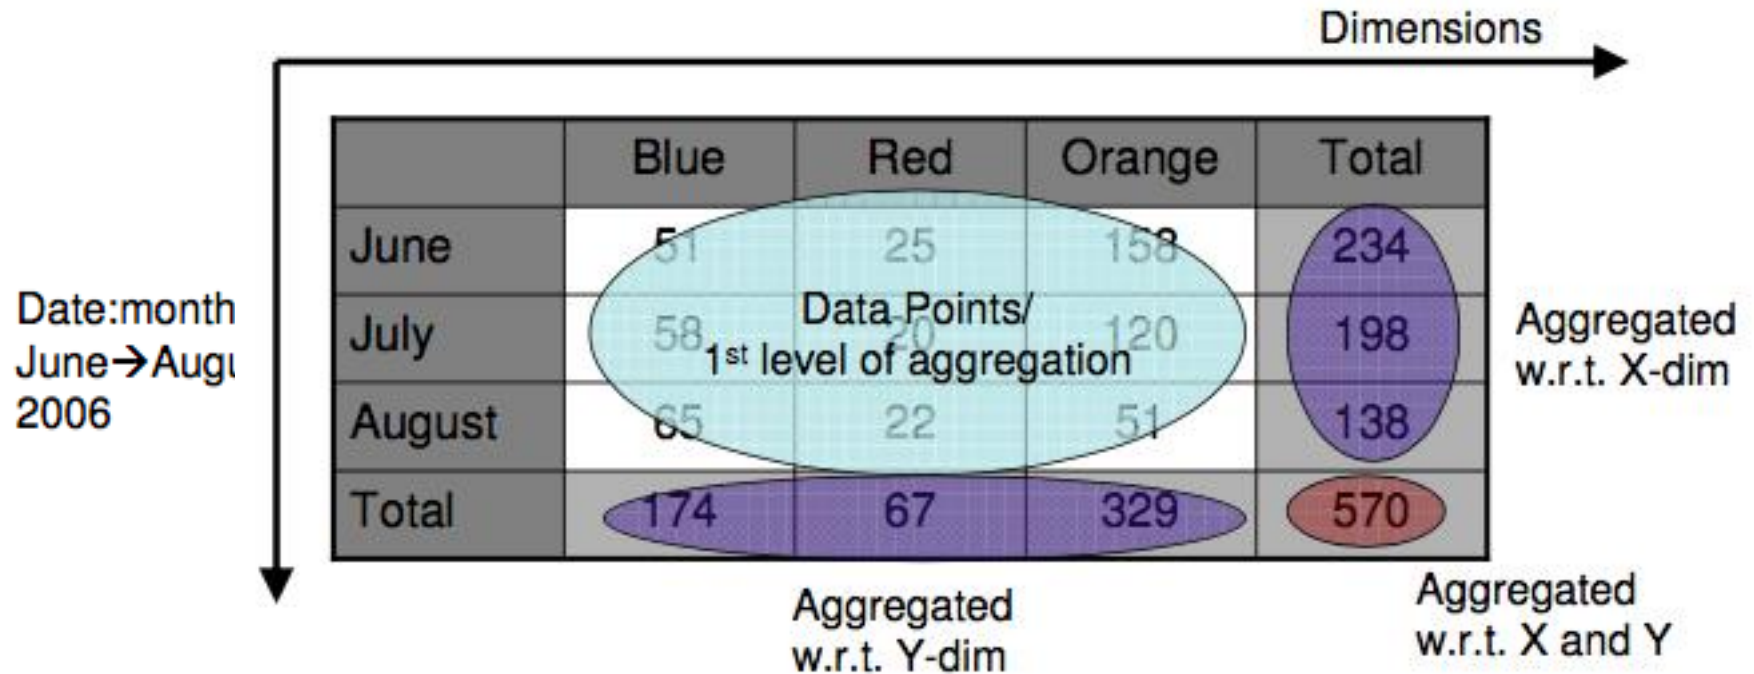



Snowflake

Schema



2 DIMENSIONAL CASE





TYPICAL OLAP OPERATIONS

What does OLAP stand for?

What OLTP?

Roll up (drill-up): summarize data

by climbing up hierarchy or by dimension reduction

Drill down (roll down): reverse of roll-up

from higher level summary to lower level summary or detailed data, or introducing new dimensions

Slice and dice: *project and select*


Pivot (rotate): *reorient the cube, visualization, 3D to series of 2D planes.*

Other operations

drill across: involving (across) more than one fact table

drill through: through the bottom level of the cube to its back-end relational tables (using SQL)

ROLLUP



REPRESENTING A CUBE

```
select semester as date, country, sum(sales)
  from sales
 group by cube(semester, country)
```

Date	Country	Sales
1st semester	Ireland	20
1st semester	France	126
1st semester	Germany	56
1st semester	null	202
2nd semester	Ireland	23
2nd semester	France	138
2nd semester	Germany	48
2nd semester	null	209
null	Ireland	43
null	France	264
null	Germany	104
null	null	411

REPRESENTING A CUBE

```
select semester as date, country, sum(sales)
  from sales
 group by roll-up(semester, country)
```

Date	Country	Sales
1st semester	Ireland	20
1st semester	France	126
1st semester	Germany	56
2nd semester	Ireland	23
2nd semester	France	138
2nd semester	Germany	48
1st semester	null	202
2nd semester	null	209
null	null	411

MDX

Multidimensional Expressions (MDX) is a query language for cubes

- Supported by many data warehouses
- Input and output are cubes

```
SELECT { [Measures].[Store Sales] } ON COLUMNS, {  
[Date].[2002], [Date].[2003] }  
ON ROWS FROM Sales  
WHERE ( [Store].[USA].[CA] )
```



ENTITY RESOLUTION

CS1951A INTRO TO DATA
SCIENCE

DATA WAREHOUSE

Application

Data Warehouse

SQL


Data Mart

SQL

CSV



DATA WAREHOUSE



DATA INTEGRATION

- **Extract-Transform-Load**
 - “Old” term
 - Schema-centric
 - Batch-oriented
- **Data Wrangling**
 - Hipster term
 - Less structured
 - Ad-hoc

DATA INTEGRATION



SCHEMA MATCHING

ID	Name	Address	Zip	State	City	Phone	E-Mail
1	Tim Kraska	135 Watermann St,	02906	Providence	RI	+1 234234 234	Tim_kraska@brown.edu
...

ID	Name	Address	Phone	E-Mail
1	Tim Kraska	135 Watermann St, 02906 Providence, RI	+1 234234 234	Tim_kraska@brown.edu
...

ID	Name
1	Tim Kraska
...	...

AddressID	Person-ID	Address	Phone-Nb	E-Mail
1	1	135 Watermann St, 02906 Providence	+1 234234 234	Tim_kraska@brown.edu
1	1	222 Hope St, 02906 Providence	980 – 0803284	tim_kraska@brown.edu

CLICKER:

HOW MANY TABLES HAS A (TYPICAL)
SAP'S ERP INSTALLATION?

- (a) 100 - 1.000
- (b) 1.000 - 10.000
- (c) 10.000 - 100.000
- (d) > 100.000

SCHEMA MATCHING: IDEAS?

ID	Name	Address	Zip	State	City	Phone	E-Mail
1	Tim Kraska	135 Watermann St,	02906	Providence	RI	+1 234234 234	Tim.kraska@brown.edu
...

ID	Name	Addr	Mobile	E-Mail
1	Tim Kraska	135 Watermann St, 02906 Providence, RI	+1 234234 234	Tim.kraska@brown.edu
...


ID	Name
1	Tim Kraska
...	...

Address ID	Person-ID	Address	Phone-Nb	E-Mail
1	1	135 Watermann St, 02906 Providence	+1 234234 234	Tim.kraska@brown.edu
1	1	222 Hope St, 02906 Providence	980 – 0803284	tim.kraska@brown.edu

SCHEMA MATCHING - TECHNIQUES

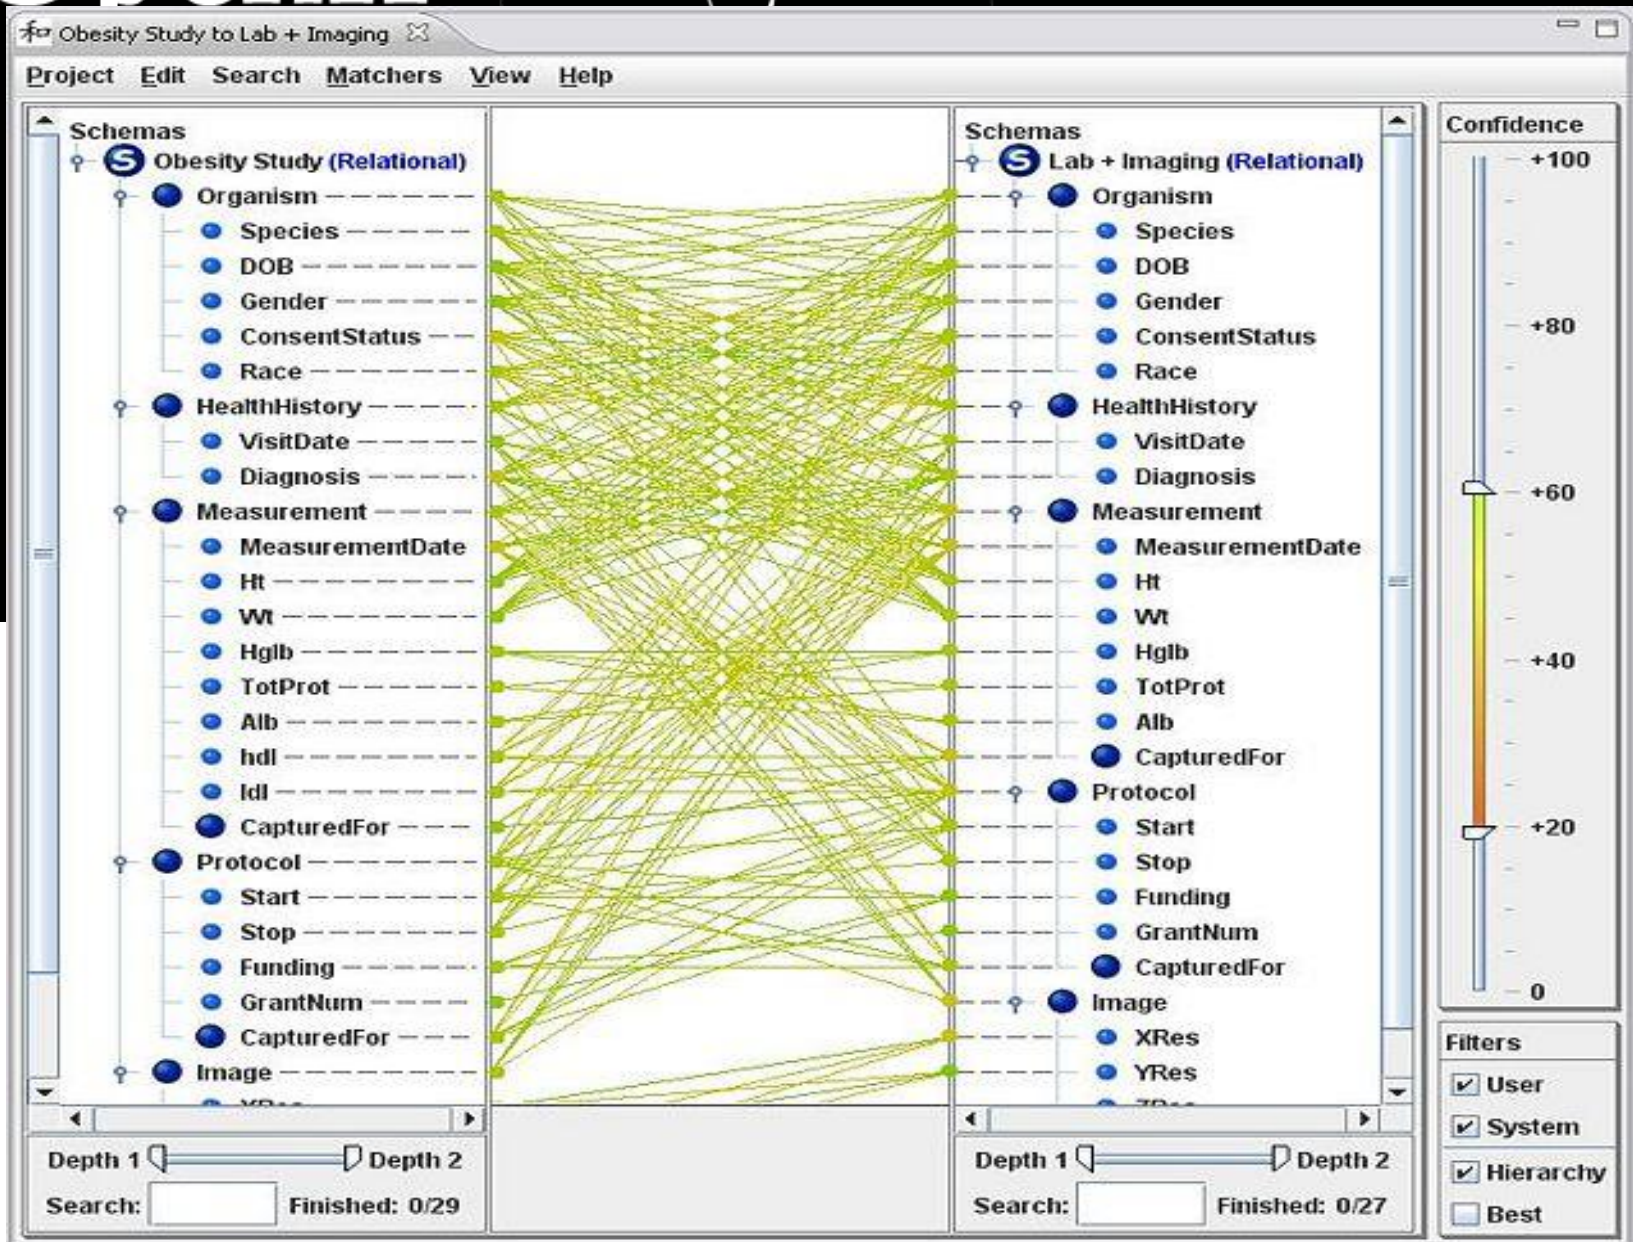
- **Instance vs Schema:** consider instance data or schema information.
- **Element vs Structure:** matching performed for individual schema element (attribute), or for combinations of elements (structure).
- **Language vs Constraint:** use linguistic information (names and textual description) or constraint information (key, relationship)
- **Matching Cardinality:** the overall match result may relate one or more elements of one schema to one or more elements of the other (1:1, 1:n, m:n).
- **Auxiliary Information:** the use of auxiliary information (dictionaries, pervious matching results, user input,..)
- Semantic difference are especially problematic (e.g., Price in Euro vs Dollar, Fahrenheit vs Celsius, etc.)

OpenII



Open Source Tools for Faster Data Integration

OpenII



DATA INTEGRATION



Schema Alignment

deduplication, entity clustering, merge/purge, fuzzy match, record linkage, approximate match...

EXAMPLE

ID	Product Name	Price
r1	iPad Two 16GB WiFi White	\$490
r2	iPad 2nd generation 16GB WiFi White	\$469
r3	iPhone 4th generation White 16GB	\$545
r4	Apple iPhone 4 16GB White	\$520
r5	Apple iPhone 3rd generation Black 16GB	\$375
r6	iPhone 4 32GB White	\$599
r7	Apple iPad2 16GB WiFi White	\$499
r8	Apple iPod shuffle 2GB Blue	\$49
r9	Apple iPod shuffle USB Cable	\$19

REAL WORLD DATA

What is wrong here?

Id	Name	Street	City	State	P-Code	Age
1	J Smith	123 UniversityAve	Seattle	Washington	98106	42
2	Mary Jones	245 3rd St	Redmond	WA	98052-1234	30
3	BobWilson	345 Broadway	Seattle	Washington	98101	19
4	M Jones	245Third Street	Redmond	NULL	98052	299
5	Robert Wilson	345 Broadway St	Seattle	WA	98101	19
6	James Smith	123 UnivAve	Seattle	WA	NULL	41
7	JWidom	123 UniversityAve	Palo Alto	CA	94305	NULL
...

REAL WORLD DATA

Customer

Inconsistent representation

Duplicate Records

Id	Name	Street	City	State	P-Code	Age
1	J Smith	123 University Ave	Seattle	Washington	98106	42
2	Mary Jones	245 3rd St	Redmond	WA	98052-1234	30
3	Bob Wilson	345 Broadway	Seattle	Washington	98101	19
4	M Jones	245 Third Street	Redmond	NULL	98052	299
5	Robert Wilson	345 Broadway St	Seattle	WA	98101	19
6	James Smith	123 Univ Ave	Seattle	WA	NULL	41
7	J Widom	123 University Ave	Palo Alto	CA	94305	NULL
...

Typos

Missing Information

REAL WORLD DATA

- How many customers do I have?

```
select count(*)  
from customer
```

Wrong answer because of duplicate records!

- How many customers by state?

```
select count(*)  
from customer  
group by state
```

State	Count
AL	60
...	...
...	...
WA	1200
Washington	50
Wasington	2

What about if you give this data to a ML algorithm?