

# SQL

INTRODUCTION TO DATA SCIENCE

TIM KRASKA



# THE 1<sup>ST</sup> PROJECT

- Goal: get started to develop a data product
- Projects can range from gathering and cleaning a data set, over repeating an existing study or web-site to visualize a data set to testing an entirely new hypothesis.
- Groups of 4
  - Piazza helps you to find group members (see special post)
- Mid-term report has to be a public blog post (afterwards you have to write weekly progress reports for the 2<sup>nd</sup> project)
- 2<sup>nd</sup> project can be a continuations of the 1<sup>st</sup>, but doesn't have to be
- 10% of your grade (the final project will be 25%).
- Today, we will publish a list of potential data sources and some project ideas on the web-page
- However, **you are encouraged to find your own data and create your own project**
- Don't worry too much about that you do not know all the tools yet

2/16/2016	Pre-Proposal Handin
3/3/2016	Advisor Checkin
3/22/2016	Mid-term report
4/12/2016	Full Project Proposal
4/21/2016	Advisor Checkin
5/3/2016	Advisor Checkin
5/12/2016	Final Project Due

# FORMAL DEFINITION OF REL. ALGEBRA

- **Atoms** (basic expressions)
- A **relation** in the database
- A **constant relation**
- **Operators** (composite expressions)
- **Selection:**  $\sigma$  (E1)
- **Projection:**  $\Pi$  (E1)
- **Cartesian Product:**  $E1 \times E2$
- **Rename:**  $\rho_V(E1)$ ,  $\rho_{A \leftarrow B}(E1)$
- **Union:**  $E1 \cup E2$
- **Minus:**  $E1 - E2$

# CODD'S THEOREM

## 3 Languages:

- **Relational Algebra**
- **Tuple Relational Calculus** (safe expressions only)
- **Domain Relational Calculus** (safe expressions only)

are **equivalent**.

## Impact of Codd's theorem:

- SQL is based on the **relational calculus**
- SQL implementation is based on **relational algebra**
- **Codd's theorem shows that SQL implementation is correct and complete.**

# NOT COVERED

**Set Division**

**Aggregate Functions**

**Codd's Proof**

**...**

# IN CLASS TASK

Player

PlayerID	Name	Age	Team
1	Russel	27	Seahawks
...	...	...	...

Team

Team	State
Seahawks	Washington
...	...

Played

PlayerID	Date	Place	Score
1	2/1/15	Phoenix	3
...	...	...	...

In relational algebra:

- 1) Return all teams, who played at least once in Phoenix
- 2) Return all Seahawks player, who did not play in the entire season

# CLICKER

Player

PlayerID	Name	Age	Team
1	Russel	27	Seahawks
...	...	...	...

Team

Team	State
Seahawks	Washington
...	...

Played

PlayerID	Date	Place	Score
1	2/1/15	Phoenix	3
...	...	...	...

Return all teams, who played at least once in Phoenix

A)  $\Pi_{\text{Team}} (\sigma_{\text{Place}=\text{'Phoenix'}} (\text{Player} \times \text{Played}))$

B)  $\Pi_{\text{Team}} \text{ Player } \bowtie (\sigma_{\text{Place}=\text{'Phoenix'}} (\text{Played}))$

C)  $\Pi_{\text{Team}} (\sigma_{\text{Place}=\text{'Phoenix'}} (\text{Player} \bowtie \text{Played}))$

# CLICKER

1.  $\Pi_{\text{Team}} (\sigma_{\text{Place}=\text{'Phoenix'}} (\text{Player} \bowtie \text{Played}))$
2.  $\Pi_{\text{Team}} \text{Player} \bowtie (\sigma_{\text{Place}=\text{'Phoenix'}} \text{Played})$
3.  $\Pi_{\text{Team}} (\sigma_{\text{Place}=\text{'Phoenix'}} (\text{Player} \bowtie \text{Played} \bowtie \text{Team}))$

**Which of these expressions are equivalent?**

- A) All
- B) 1 and 2
- C) 1 and 3
- D) 2 and 3
- E) None



# CLICKER

Player

PlayerID	Name	Age	Team
1	Russel	27	Seahawks
...	...	...	...

Team

Team	State
Seahawks	Washington
...	...

Played

PlayerID	Date	Place	Score
1	2/1/15	Phoenix	3
...	...	...	...

Return all Seahawks player names, who did not play so far

- A)  $\Pi_{\text{Name}} ( \sigma_{\text{Team}=\text{'Seahawks'}} (\text{Player} \bowtie \text{Played}))$
- B)  $\Pi_{\text{Name}} ( \sigma_{\text{Team}=\text{'Seahawks'}} (\text{Player})) - \Pi_{\text{Name}} ( \sigma_{\text{Team}=\text{'Seahawks'}} (\text{Player} \bowtie \text{Played}))$
- C)  $\Pi_{\text{Name}} ( \sigma_{\text{Team}=\text{'Seahawks'}} (\text{Player} \times \text{Played}))$
- D)  $\Pi_{\text{Name}} ( \sigma_{\text{Team}=\text{'Seahawks'}} \wedge \text{Date} = \text{null}) (\text{Player} \bowtie \text{Played}))$

# (SIMPLE) DATA DEFINITION WITH SQL

## Data types:

- character (n), char (n)
- character varying (n), varchar (n)
- numeric (p,s), integer
- blob or raw for large binaries
- clob for large string values

## Create Tables:

```
create table Professor
  (Person-ID   integer not null,
   Name        varchar (30) not null
   Level       varchar (2) default AP);
```

# DDL (CTD.)

## **Delete a Table:**

```
drop table Professor;
```

## **Modify the structure of a Table:**

```
alter table Professor add column(age integer);
```

## **Management of Indexes (Performance tuning):**

```
create index myIndex on Professor(name, age);  
drop index myIndex;
```


# UPDATES (DML)

## Insert Tuples

```
insert into Student (Student-ID, Name)  
  values (28121, `Archimedes`);
```

```
insert into attends  
  select Student-ID, Course-ID  
  from Student, Lecture  
  where Title= `Logic` ;
```

Student		
Student-ID	Name	Semester
29120	Theophrastos	2
29555	Feuerbach	2
28121	Archimedes	-



# SEQUENCE TYPES (AUTOMATIC INCREMENT FOR SURROGATES)

```
create sequence Person-ID_seq increment by 1 start with 1;  
insert into Professor(Person-ID, Name)  
  values(Person-ID_seq.nextval, „Roscoe“);
```

Syntax is vendor dependent

E.g., AUTO-INCREMENT Option in MySQL

Syntax above was standardized in SQL 2003

# UPDATES (CTD.)

## Delete tuples

```
delete Student  
where Semester > 13;
```

## Update tuples

```
update Student  
    set Semester = Semester + 1;
```

# QUERIES

```
select      Person-ID, Name
from        Professor
where       Level = 'FP';
```

Person-ID	Level	Name	Room
2125	FP	Ugur	303
2126	FP	Stan	345
2165	AP	Tim	335
2136	FP	Curie	401
2137	AP	Jeff	507



Person-ID	Name
2125	Ugur
2126	Stan
2136	Curie

$\Pi_{\text{Person-ID, Name}} (\sigma_{\text{Level}='FP'} (\text{Professor}))$



# QUERIES: SORTING

```
select Person-ID, Name, Level  
from Professor  
order by Level desc, Name asc;
```

Person-ID	Name	Level
2136	Curie	FP
2137	Jeff	FP
2126	Stan	FP
2125	Ugur	FP
2134	Augustinus	AP
2127	Kopernikus	AP
2133	Popper	AP

# CLICKER QUESTION: ARE THE FOLLOWING QUERIES EQUIVALENT?

**select** Level  
**from** Professor

$\Pi_{\text{Level}}$  (Professor)

Answer:

(1) Yes

(2) No

Person-ID	Name	Level
2136	Curie	FP
2137	Jeff	FP
2126	Stan	FP
2125	Ugur	FP
2134	Augustinus	AP
2127	Kopernikus	AP
2133	Popper	AP

# DUPLICATE ELIMINATION

**select distinct** Level  
**from** Professor

Level
AP
FP

# QUERIES: JOINS

## Who teaches ML?

```
select Name  
from Professor, Lecture  
where Person-ID = ProfID and Title = `ML` ;
```

$$\Pi_{\text{Name}}(\sigma_{\text{Person-ID}=\text{Prof-ID} \wedge \text{Title}=\text{'ML'}}(\text{Professor} \times \text{Lecture}))$$

Renamed Lecture.Person-ID to Prof-ID

Will show later how this can be done as part of a query.

# JOINS

Person-ID	Name	Level	Room
2165	Ugur	FP	226
2166	Stan	FP	232
...	...	...	...
2200	Jeff	FP	7

CID	Title	CP	Prof-ID
5001	Foundation	4	2137
5041	German	4	2125
...	...	...	...
5049	ML	2	2125
...	...	...	...
4630	Vision	4	2137

Professor x Lecture

PID	Name	Level	Room	CID	Title	CP	ProfID
2125	Ugur	FP	226	5001	Foundation	4	2137
1225	Ugur	FP	226	5041	German	4	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

PID	Name	Level	Room	CID	Title	CP	ProfID
2125	Ugur	FP	226	5001	Foundation	4	2137
1225	Ugur	FP	226	5041	German	4	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2125	Ugur	FP	226	5049	ML	2	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2126	Stan	FP	232	5001	ML	4	2137
2126	Stan	FP	232	5041	German	4	2125
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2137	Jeff	FP	7	4630	Vision	4	2137

↓  $\sigma_{\text{Person-ID}=\text{Prof-ID} \wedge \text{Title}=\text{'ML'}}$

Person-ID	Name	Level	Room	ID	Title	CP	ProfID
2125	Ugur	FP	226	5049	ML	2	2125

↓  $\Pi_{\text{Name}}$


Name
Ugur

# SQL -> RELATIONAL ALGEBRA

## SQL

**select**  $A_1, \dots, A_n$   
**from**  $R_1, \dots, R_k$   
**where**  $P$ ;

## Relational Algebra



# WHO ATTENDS WHICH LECTURE?

Professor(Person-ID:integer, Name:string)


Student(Student-ID:integer, Name:string)

Lecture(Course-ID:string, Title:string, CP:float)

Gives(Person-ID:integer, Course-ID:string)

Attends(Student-ID:integer, Course-ID:string)

Tests(Student-ID:integer, Course-ID:string, Person-ID:integer, ,  
Grade:string)



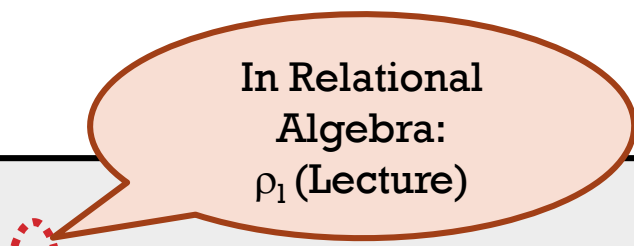


# JOINS AND TUPLE VARIABLES

**Equivalent queries: Who attends which lecture?**

```
select Name, Title  
from Student, attends, Lecture  
where Student.Student-ID = attends.Student-ID and  
        attends.Course-ID = Lecture. Course-ID;
```

```
select s.Name, l.Title  
from Student s, attends a, Lecture l  
where s.Student-ID = a.Student-ID and  
        a.Course-ID = l.Course-ID;
```



In Relational  
Algebra:  
 $\rho_1$  (Lecture)

# RENAME OF ATTRIBUTES

**Give title and professor of all lectures?**

```
select Title, Person-ID as ProfID  
from Lecture;
```

# SET OPERATIONS

## Union, Intersect, Minus

```
( select Name  
  from Assistant )  
union  
( select Name  
  from Professor);
```

# **GROUPING AGGREGATION**

# GROUPING, AGGREGATION

**Aggregate functions: avg, max, min, count, sum**

```
select avg (Semester)
  from Student;
```

```
select Person-ID, sum (CP) as load
  from Lecture
 group by Person-ID;
```

```
select p.Person-ID, Name, sum (CP)
  from Lecture l, Professor p
 where l.Person-ID= p.Person-ID and level = 'FP'
 group by p.Person-ID, Name
 having avg (CP) >= 3;
```

# IMPERATIVE PROCESSING IN SQL

Step 1:

**from** Lecture l, Professor p

**where** l.Person-ID= p.Person-ID **and** level = 'FP

Step 2:

**group by** p.Person-ID, Name

Step 3:


**having avg** (CP)  $\geq 3$ ;

Step 4:

**select** p.Person-ID, Name, **sum** (CP)


# GROUP BY

Lecture x Professor							
Nr	Title	CP	Person-ID	Person-ID	Name	Level	Room
5001	Foundation	4	2137	2125	Ugur	FP	226
5041	German	4	2125	2125	Ugur	FP	226
...	...	...	...	...	...	...	...
4630	Vision	4	2137	2137	Jeff	AP	7



(where l.Person-ID= p.Person-ID and level = 'FP')

Nr	Title	CP	Person-ID	Person-ID	Name	Level	Room
5001	Foundation	4	2137	2137	Jeff	FP	7
5041	German	4	2125	2125	Ugur	FP	226
5043	Cyper Stuff	3	2126	2126	Stan	FP	232
5049	ML	2	2125	2125	Ugur	FP	226
4052	Logik	4	2125	2125	Ugur	FP	226
5052	Robotics	3	2126	2126	Stan	FP	232
5216	Adv. German	2	2126	2126	Stan	FP	232
4630	Vision	4	2137	2137	Jeff	FP	7



**group by** p.Person-ID, Name





**group by** p.Person-ID, Name

Nr	Title	CP	Person-ID	Person-ID	Name	Level	Room
5041	German	4	2125	2125	Ugur	FP	226
5049	ML	2	2125	2125	Ugur	FP	226
4052	Logik	4	2125	2125	Ugur	FP	226
5043	Cyper Stuff	3	2126	2126	Stan	FP	232
5052	Robotics.	3	2126	2126	Stan	FP	232
5216	Adv. German	2	2126	2126	Stan	FP	232
5001	Foundation	4	2137	2137	Jeff	FP	7
4630	Vision	4	2137	2137	Jeff	FP	7



**Having** (CP) >= 3

Nr	Title	CP	Person-ID	Person-ID	Name	Level	Room
5041	German	4	2125	2125	Ugur	FP	226
5049	ML	2	2125	2125	Ugur	FP	226
4052	Logik	4	2125	2125	Ugur	FP	226
5001	Foundation	4	2137	2137	Jeff	FP	7
4630	Vision	4	2137	2137	Jeff	FP	7

Nr	Title	CP	Person-ID	Person-ID	Name	Level	Room
5041	German	4	2125	2125	Ugur	FP	226
5049	ML	2	2125	2125	Ugur	FP	226
4052	Logik	4	2125	2125	Ugur	FP	226
5001	Foundation	4	2137	2137	Jeff	FP	7
4630	Vision	4	2137	2137	Jeff	FP	7


















**sum(select)**

Person-ID	Name	sum (CP)
2125	Ugur	10
2137	Jeff	8

Question: Why do we need to group-by on Person-ID and Name?

# CLICKER

Which of the following is the correct order of keywords for SQL SELECT statements?

- A) From, Where, Select, Group-By, Having
- B) Select, From, Where, Group-by, Having
- C) Having, Select, From, Where, Group-by
- D) From, Where, Group-By, Having, Select

What is the execution order?

- A) From, Where, Select, Group-By, Having
- B) Select, From, Where, Group-by, Having
- C) Having, Select, From, Where, Group-by
- D) From, Where, Group-By, Having, Select

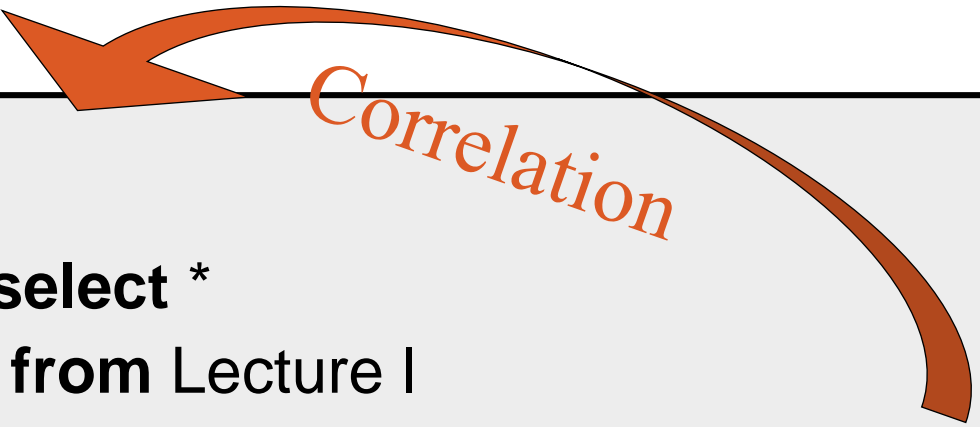
# SUBQUERIES

# EXISTENTIAL QUANTIFICATION

## EXISTS SUBQUERIES

```
select p.Name  
from Professor p  
where not exists ( select *  
                    from Lecture l  
                    where l.Person-ID = p.Person-ID );
```

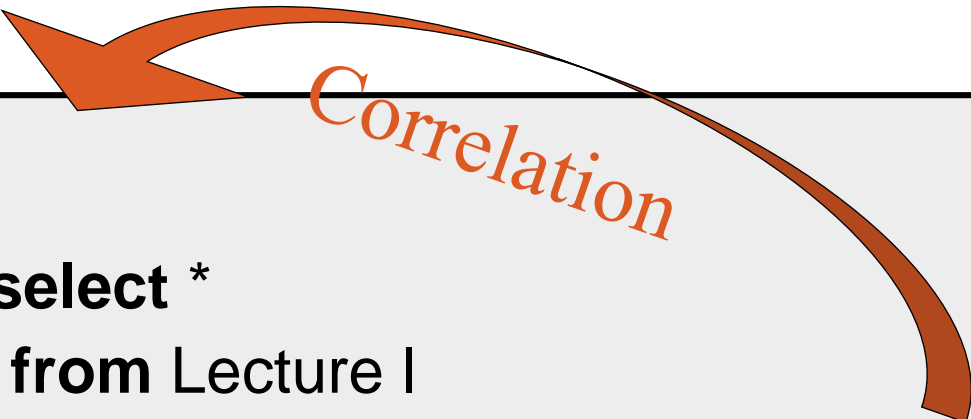
# CORRELATED SUB-QUERIES



The diagram illustrates the concept of correlation in a SQL query. A large, curved orange arrow originates from the word 'Correlation' and points to the 'p' in the subquery's 'where' clause. The word 'Correlation' is written in a large, orange, italicized font, following the curve of the arrow. The SQL query is displayed in a light gray box with a black border. The query is:   
**select** p.Name   
**from** Professor p   
**where not exists** ( **select** \*   
                    **from** Lecture l   
                    **where** l.Person-ID = p.Person-ID );

```
select p.Name
from Professor p
where not exists ( select *
                   from Lecture l
                   where l.Person-ID = p.Person-ID );
```

# CORRELATED SUB-QUERIES



```
select p.Name
from Professor p
where not exists ( select *
                   from Lecture l
                   where l.Person-ID = p.Person-ID );
```

For every p in Professor  
 where not exist ( // like a check for empty set  
 select \* from lecture where l.Person-ID = p.Person-ID  
 )  
 emit p.Name

# UNCORRELATED SUB-QUERY

```
select Name  
from Professor  
where Person-ID not in ( select Person-ID  
                           from Lecture);
```

What is better? Correlated or uncorrelated?



# SUB-QUERIES WITH ALL

```
select Name  
from Student  
where Semester >= all ( select Semester  
                        from Student);
```

# SUBQUERIES IN SELECT, FROM

```
select Person-ID, Name, ( select sum (CP) as load  
                           from Lecture l  
                           where p.Person-ID=l.Person-ID )  
from Professor p;
```

```
select p.Person-ID, Name, pl.load  
from Professor p, ( select Person-ID, sum (CP) as load  
                   from Lecture  
                   group by Person-ID ) pl  
where p.Person-ID = pl.Person-ID;
```

Is this better than the simple Group By Query  
from before?

# QUERY REWRITE

**Two equivalent join queries: Which is better?**

```
select *  
from Assistant a  
where exists  
    (select *  
     from Professor p  
     where a.Boss = p.Person-ID and p.age < a.age);
```

```
select a.*  
from Assistant a, Professor p  
where a.Boss=p.Person-ID and p.age < a.age;
```

# ARE THESE TWO QUERIES EQUIVALENT?

```
select count (*)  
from Student  
where Semester < 13 or  
Semester > =13;
```

```
select count (*)  
from Student;
```

**(A) No      (B) Yes**

# NULL VALUES

# NULL VALUES (NULL = UNKNOWN)

**Are these two queries equivalent?**

```
select count (*)  
from Student  
where Semester < 13 or  
Semester > =13;
```

```
select count (*)  
from Student;
```

# WORKING WITH NULL VALUES

**Arithmetics: Propagate null:** If an operand is null, the result is null.

- $\text{null} + 1 \rightarrow \text{null}$
- $\text{null} * 0 \rightarrow \text{null}$

**Comparisons: New Boolean value unknown.** All comparisons that involve a null value, evaluate to unknown.

- $\text{null} = \text{null} \rightarrow \text{unknown}$
- $\text{null} < 13 \rightarrow \text{unknown}$
- $\text{null} > \text{null} \rightarrow \text{unknown}$

**Logic: Boolean operators are evaluated using the following tables (next slide):**

# LOGICAL OPERATIONS

<b>p</b>	<b>NOT p</b>
TRUE	FALSE
FALSE	TRUE
Unknown	Unknown

<b>p</b>	<b>q</b>	<b>p OR q</b>	<b>p AND q</b>	<b>p = q</b>
TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	FALSE	TRUE	FALSE	FALSE
TRUE	Unknown	TRUE	Unknown	Unknown
FALSE	TRUE	TRUE	FALSE	FALSE
FALSE	FALSE	FALSE	FALSE	TRUE
FALSE	Unknown	Unknown	FALSE	Unknown
Unknown	TRUE	TRUE	Unknown	Unknown
Unknown	FALSE	Unknown	FALSE	Unknown
Unknown	Unknown	Unknown	Unknown	Unknown



**where:** Only tuples which evaluate to **true** are part of the query result. (**unknown** and **false** are equivalent here):

```
select count (*)  
from Student  
where Semester < 13 or Semester > =13;
```

**group by:** If exists, then there is a group for **null**.

```
select count (*)  
from Student  
group by Semester;
```

Predicates with null:

```
select count (*) from Student  
where Semester is null;
```

# CLICKER

SELECT count(\*) AS total FROM orders;

Count(*)
100

SELECT count(\*) AS cust\_123\_total FROM orders  
WHERE customer\_id = '123';

Count(*)
15

**Given the above query results, what will be the result of the query below?**

SELECT count(\*) AS cust\_not\_123\_total  
FROM orders  
WHERE customer\_id <> '123'

A) 85    B) 100    C) Impossible to say

# SYNTACTIC SUGAR

```
select *  
from Student  
where Semester > = 1 and Semester < = 6;
```

```
select *  
from Student  
where Semester between 1 and 6;
```

```
select *  
from Student  
where Semester in (2,4,6);
```

# COMPARISONS WITH LIKE

"%,, represents any sequence of characters (0 to n)

"\_,, represents exactly one character

N.B.: For comparisons with = , % and \_ are normal chars.

```
select *
```

```
from Student
```

```
where Name like 'Tim%';
```

```
select distinct Name
```

```
from Lecture l, attends a, Student s
```

```
where s.Student-ID = a.Student-ID and a.CID = l.CID  
      and l.Title like '%science%';
```

# JOINS IN SQL-92

- **cross join**: Cartesian product
- **natural join**
- **join** or **inner join**
- **left, right** or **full outer join**
- (union join: not discussed here)

```
select *  
from R1, R2  
where R1.A = R2.B;
```

```
select *  
from R1 join R2 on R1.A = R2.B;
```

# LEFT OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID,  
s.Student-ID, s.Name  
from Professor p left outer join  
    (tests t left outer join Student s  
    on t.Student-ID= s.Student-ID)  
    on p.Person-ID=t.Person-ID;
```

Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
2136	Curie	-	-	-	-	-
⋮	⋮	⋮	⋮	⋮	⋮	⋮

# RIGHT OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID, s.Student-ID,
s.Name
from Professor p right outer join
      (tests t right outer join Student s on
      t.Student-ID= s.Student-ID)
on p.Person-ID=t.Person-ID;
```

Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
-	-	-	-	-	26120	Fichte
⋮	⋮	⋮	⋮	⋮	⋮	⋮

# FULL OUTER JOINS

```
select p.Person-ID, p.Name, t.Person-ID, t.Grade, t.Student-ID,  
s.Student-ID, s.Name
```

```
from Professor p full outer join
```

```
    (tests t full outer join Student s on  
    t.Student-ID= s.Student-ID)
```

```
on p.Person-ID=t.Person-ID;
```

p.Person-ID	p.Name	t.Person-ID	t.Grade	t.Student-ID	s.Student-ID	s.Name
2126	Stan	2126	1	28106	28106	Carnap
2125	Ugur	2125	2	25403	25403	Jonas
2137	Jeff	2137	2	27550	27550	Schopenhauer
-	-	-	-	-	26120	Fichte
2136	Curie	-	-	-	-	-